

Housing Prices Project

Sulayman Desai

Abstract

This project required me to predict sale prices from the features of a housing dataset. I used a number of data processing techniques such as imputing missing values, encoding categorical variables, normalising features and removing outliers to clean up the training and testing data sets. Using these cleaned data sets, I fitted various machine learning models such as: Ordinary Least Squares Regression, Ridge Regression, Lasso Regression, Elastic Net Regression, Decision Tree Regressor, Random Forest Regressor and XGBoost Regressor. The **Ridge Regression** model returned the lowest **mean absolute error** of **0.07871660662991475**. It was found that the tree models demanded great computational cost for minimal improvements in accuracy. If this project were to be improved, I would implement feature generation to generate more useful features, used different encoding techniques for greater usage of categorical variables and would develop an ensemble of models to get the best predictions.

Contents

1	Housing Prices Competition	3
2	Importing Libraries	3
3	Importing Data	4
4	Examining the Datasets	4
5	Processing Data	11
5.0.1	Checking skewness and kurtosis	11
5.0.2	Sale Price log-transformation	13
5.0.3	Missing Values	14
5.0.4	Imputation	18
5.0.5	Dealing with Outliers	19
5.0.6	Encoding Categorical Variables	23
5.0.7	Checking skewness and kurtosis of rest of dataset	25
5.0.8	And beyond...	29
6	Cleaning Test Data	31
6.0.1	Missing Values	32
6.0.2	Outliers	38
6.0.3	Checking skewness	42
6.0.4	Categorical Variables	44
6.0.5	Comparing Test and Train datasets	46
7	Modelling	58
7.0.1	Linear Regression	58
7.0.2	Ridge Regression	61
7.0.3	Lasso Regression	62
7.0.4	Elastic Net	63
7.0.5	Decision Tree Regressor	64
7.0.6	Random Forest Regressor	66
7.0.7	XGBoost Regressor	67
8	Conclusion	68

1 Housing Prices Competition

In this notebook, I will be using the Housing Prices dataset (<https://www.kaggle.com/c/house-prices-advanced-regression-techniques/data>) to predict future housing prices by applying different linear regression and decision tree models.

2 Importing Libraries

I am importing the relevant libraries which will be used.

```
[1]: # Importing packages to deal with dataframes and numerical functions
import pandas as pd
import numpy as np
from scipy import stats

# Importing modules for graphing data
import matplotlib.pyplot as plt
import matplotlib as mpl
import seaborn as sns

# stores plot within notebook
%matplotlib inline

# Importing package which makes plots have same theme as notebook, using jupyter.
→style()
#from jupyterthemes import jtplot

# corrects behaviour between panda and matplotlib libraries
pd.plotting.register_matplotlib_converters()

# Python Machine Learning (ML) library Scikit-Learn
import sklearn
import xgboost
import category_encoders as ce

# Importing ML models
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge, Lasso
from sklearn.ensemble import GradientBoostingRegressor, RandomForestRegressor
from sklearn.tree import DecisionTreeRegressor
from xgboost import XGBRegressor

# Importing error metrics
from sklearn.metrics import mean_squared_error, mean_absolute_error
from scipy.stats import kurtosis, skew

# Importing data manipulation tools
```

```

from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_validate, cross_val_score, \
    ↳StratifiedKFold
from sklearn.model_selection import KFold
from sklearn.preprocessing import StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import LabelEncoder, OneHotEncoder, LabelBinarizer, \
    ↳MultiLabelBinarizer
from sklearn.preprocessing import PolynomialFeatures
from sklearn import preprocessing
from scipy.stats import norm

# ensures cross platform functionality
import os

```

3 Importing Data

Now I will import the train, test and sample data.

```

[2]: # importing training data
train_df = pd.read_csv(r'train.csv')

# importing test data
test_df = pd.read_csv(r'test.csv')

# importing sample submission data
sample_df = pd.read_csv(r'sample_submission.csv')

```

4 Examining the Datasets

I am going to check the structure of the training set, discovering questions such as, how large the data is, what the data contains, are there any missing values etc.

First, let me get a quick look at the first few rows of the data.

```

[3]: # displaying the first 5 rows of data
train_df.head()

```

```

[3]:   Id  MSSubClass MSZoning  LotFrontage  LotArea Street Alley LotShape  \
0    1           60       RL           65.0     8450   Pave   NaN     Reg
1    2           20       RL           80.0     9600   Pave   NaN     Reg
2    3           60       RL           68.0    11250   Pave   NaN    IR1
3    4           70       RL           60.0     9550   Pave   NaN    IR1
4    5           60       RL           84.0    14260   Pave   NaN    IR1

```

	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal	MoSold	\
0	Lvl	AllPub	...	0	NaN	NaN	NaN	0	2	
1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	5	
2	Lvl	AllPub	...	0	NaN	NaN	NaN	0	9	
3	Lvl	AllPub	...	0	NaN	NaN	NaN	0	2	
4	Lvl	AllPub	...	0	NaN	NaN	NaN	0	12	

	YrSold	SaleType	SaleCondition	SalePrice
0	2008	WD	Normal	208500
1	2007	WD	Normal	181500
2	2008	WD	Normal	223500
3	2006	WD	Abnorml	140000
4	2008	WD	Normal	250000

[5 rows x 81 columns]

I want to get some general information about the data set which may be useful later on.

```
[4]: # displaying various statistical information about training data
train_df.describe()
```

```
[4]:
```

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	\
count	1460.000000	1460.000000	1201.000000	1460.000000	1460.000000	
mean	730.500000	56.897260	70.049958	10516.828082	6.099315	
std	421.610009	42.300571	24.284752	9981.264932	1.382997	
min	1.000000	20.000000	21.000000	1300.000000	1.000000	
25%	365.750000	20.000000	59.000000	7553.500000	5.000000	
50%	730.500000	50.000000	69.000000	9478.500000	6.000000	
75%	1095.250000	70.000000	80.000000	11601.500000	7.000000	
max	1460.000000	190.000000	313.000000	215245.000000	10.000000	

	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea	BsmtFinSF1	...	\
count	1460.000000	1460.000000	1460.000000	1452.000000	1460.000000	...	
mean	5.575342	1971.267808	1984.865753	103.685262	443.639726	...	
std	1.112799	30.202904	20.645407	181.066207	456.098091	...	
min	1.000000	1872.000000	1950.000000	0.000000	0.000000	...	
25%	5.000000	1954.000000	1967.000000	0.000000	0.000000	...	
50%	5.000000	1973.000000	1994.000000	0.000000	383.500000	...	
75%	6.000000	2000.000000	2004.000000	166.000000	712.250000	...	
max	9.000000	2010.000000	2010.000000	1600.000000	5644.000000	...	

	WoodDeckSF	OpenPorchSF	EnclosedPorch	3SsnPorch	ScreenPorch	\
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	
mean	94.244521	46.660274	21.954110	3.409589	15.060959	
std	125.338794	66.256028	61.119149	29.317331	55.757415	
min	0.000000	0.000000	0.000000	0.000000	0.000000	

25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	25.000000	0.000000	0.000000	0.000000
75%	168.000000	68.000000	0.000000	0.000000	0.000000
max	857.000000	547.000000	552.000000	508.000000	480.000000

	PoolArea	MiscVal	MoSold	YrSold	SalePrice
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000
mean	2.758904	43.489041	6.321918	2007.815753	180921.195890
std	40.177307	496.123024	2.703626	1.328095	79442.502883
min	0.000000	0.000000	1.000000	2006.000000	34900.000000
25%	0.000000	0.000000	5.000000	2007.000000	129975.000000
50%	0.000000	0.000000	6.000000	2008.000000	163000.000000
75%	0.000000	0.000000	8.000000	2009.000000	214000.000000
max	738.000000	15500.000000	12.000000	2010.000000	755000.000000

[8 rows x 38 columns]

As the dataset has a large number of columns, not all of them are showing when displaying the dataframe, so I will print all the column names. It could also be useful to know the shape of the dataframe.

```
[5]: # Printing all the column names
print(train_df.columns)
train_df.shape
```

```
Index(['Id', 'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street',
      'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig',
      'LandSlope', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType',
      'HouseStyle', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd',
      'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType',
      'MasVnrArea', 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual',
      'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1',
      'BsmtFinType2', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating',
      'HeatingQC', 'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF',
      'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath',
      'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual',
      'TotRmsAbvGrd', 'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType',
      'GarageYrBlt', 'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual',
      'GarageCond', 'PavedDrive', 'WoodDeckSF', 'OpenPorchSF',
      'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'PoolQC',
      'Fence', 'MiscFeature', 'MiscVal', 'MoSold', 'YrSold', 'SaleType',
      'SaleCondition', 'SalePrice'],
      dtype='object')
```

```
[5]: (1460, 81)
```

I want to see how varied the instances of Sale Price are in the training data, so I will plot a univariate distribution which shows the probability of occurrence of Sale Price.

```
[6]: # Setting the size of the figure
plt.figure("sale_price_histogram", figsize=(15,8))


# Giving the graph a title
plt.title("Univariate distribution of Sale Price in the training data", fontsize=14)

plt.ylabel("Frequency", fontsize = 14)

plt.xlabel("Sale Price", fontsize = 14)

# Plotting univariate distribution of sale price
sns.distplot(train_df.SalePrice, axlabel = False)

plt.show;
```



This graph clearly shows that the distribution peaks somewhere inbetween 100,000-200,000 dollars. The SalePrice has a positive skewness away from the normal distribution.

I think it would be worth to just get a general overview of how correlated all the features are with each other.


```
[7]: plt.figure('correlation_matrix_training', figsize = (15,8))

# Calculating the correlation matrix of the dataframe train_df
corr_train = train_df.corr()

sns.heatmap(data=corr_train, square = True)
```

```
plt.title('Correlation matrix of training data')

plt.show;
```



Clearly, this does not provide too much information about the dataset as it is difficult to read. Hence, I will check which features have the most impact on Sale Price by finding the top 10 correlated features.

```
[8]: # Creating a function which returns the n top correlations with SalePrice
def get_top_correlations(df, n):
    #df.corrwith returns the correlation with SalePrice
    a_corr = df.corrwith(df['SalePrice'])
    # implementing a loop through the rows of the series au_corr to pick the n
    →largest correlations
    for row in a_corr:
        n_largest = a_corr.nlargest(n)
        return n_largest
```



```
[9]: # Calling function get_top_correlations to find the top 10 correlating features
      ↳with SalePrice
      correlations_top_10 = get_top_correlations(train_df, 10)

      # Printing the correlation coefficients
      print(correlations_top_10)
```

```
SalePrice      1.000000
OverallQual    0.790982
GrLivArea      0.708624
GarageCars     0.640409
GarageArea     0.623431
TotalBsmtSF    0.613581
1stFlrSF       0.605852
FullBath       0.560664
TotRmsAbvGrd   0.533723
YearBuilt      0.522897
dtype: float64
```

```
[10]: #Plotting a bar chart to easily identify which features are most correlated with
      ↳SalePrice

      plt.figure('top_10_correlated_saleprice',figsize=(15,8))

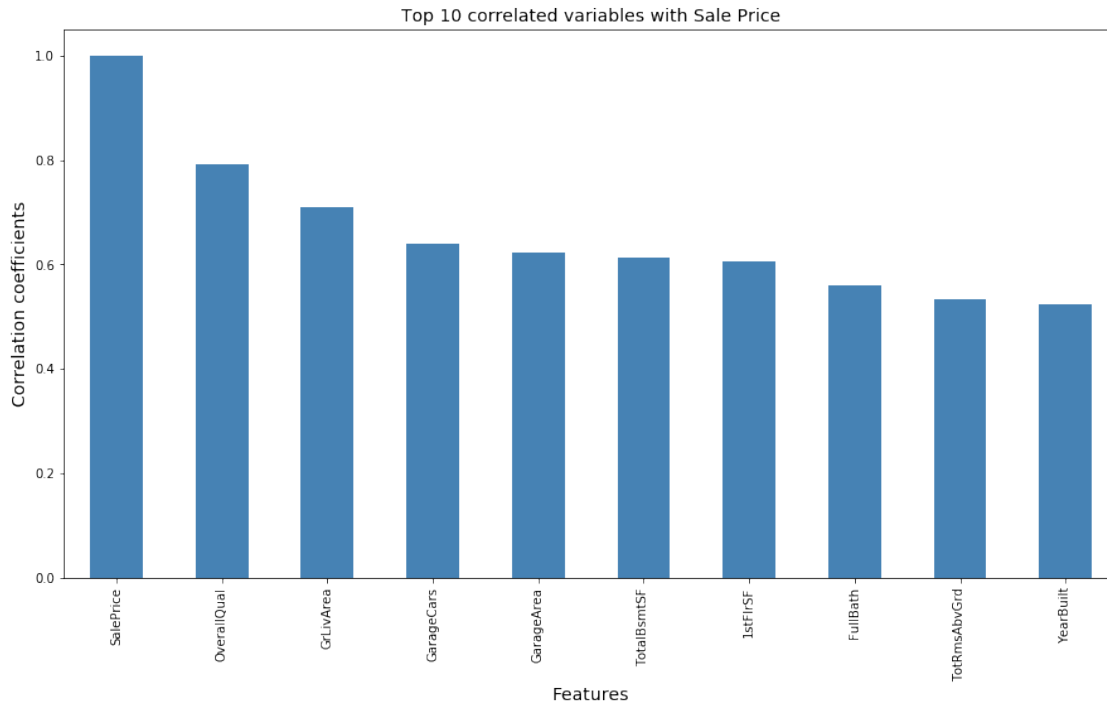
      plt.title("Top 10 correlated variables with Sale Price", fontsize = 14)

      plt.ylabel("Correlation coefficients", fontsize = 14)

      plt.xlabel("Features", fontsize = 14)

      correlations_top_10.plot.bar(color = 'steelblue')

      plt.show;
```



As seen in the barplot, OverallQual, GrLivArea and GarageCars are all strongly correlated with Sale Price. The rest of the features are moderately correlated with Sale Price.


The above graph shows the top 10 correlated variables with Sale Price. But it would be nice to see scatter graphs of the strongly correlated numerical features with Sale Price. Categorical features (for the moment) cannot demonstrate any correlation with Sale Price as they do not have any values assigned to them.

```
[11]: # Selecting only the columns which have numerical values and plotting a scatter
      ↪ plot against SalePrice
top_correlated_feat = ['SalePrice', 'OverallQual', 'GrLivArea', 'GarageCars',
      ↪ 'TotalBsmtSF', 'FullBath', 'YearBuilt']

sale_price_scatter = sns.pairplot(train_df[top_correlated_feat], y_vars =
      ↪ ['SalePrice'],
      x_vars = ['OverallQual', 'GrLivArea', 'GarageCars', 'TotalBsmtSF',
      ↪ 'FullBath', 'YearBuilt'],
      height = 2.5)

sale_price_scatter.fig.suptitle("Relationship between Sale Price and strongly
      ↪ correlated numerical features", y=1.08)

plt.show;
```



The above scatter plots give a very easy visual representation of the relationship of different numerical features with sale price.

As expected, the higher the overall quality of the housebuild, the higher the sale price. This is also the case for ground living area. We had confirmed both of these before from our correlation coefficients, but now we can visually see the relationship.

I also find interesting that houses built after 2000 increase in price a lot more than before 2000. I think it may be worth checking what is different about these houses, and possibly speaking to a housing expert.

5 Processing Data

5.0.1 Checking skewness and kurtosis

In the univariate distribution plot of SalePrice, we saw that the training data had a positive skewness. I want to check the actual value of the skewness and to check another piece of information about the SalePrice, the kurtosis. Skewness is a measure of how symmetric a distribution is, whilst kurtosis is a measure of the heaviness of the tail of the distribution i.e. a measure of the number of outliers in the data.

Skewness can be calculated using the following formula:

$$s = \frac{\sum_{i=1}^N (x_i - \bar{x})^3}{\sigma^3}$$

whilst, kurtosis can be calculated using the following formula:

$$K = \frac{1}{N} \sum_{i=1}^N \left(\frac{x_i - \mu}{\sigma} \right)^4$$

```
[12]: # Calculating skewness and kurtosis

skew_SalePrice = skew(train_df['SalePrice'])
print('Skewness of sale price is:', skew_SalePrice)

kurt_SalePrice = kurtosis(train_df['SalePrice'])
print('Kurtosis of sale price is:', kurt_SalePrice)
```

Skewness of sale price is: 1.880940746034036

Kurtosis of sale price is: 6.509812011089439

As seen in the values of skewness and kurtosis, SalePrice demonstrates it has positive skewness and a significant positive kurtosis. This means the data is not normally distributed, as normally distributed data will have a skewness of 0 and a kurtosis of 0 (Python uses Fisher's definition as default).

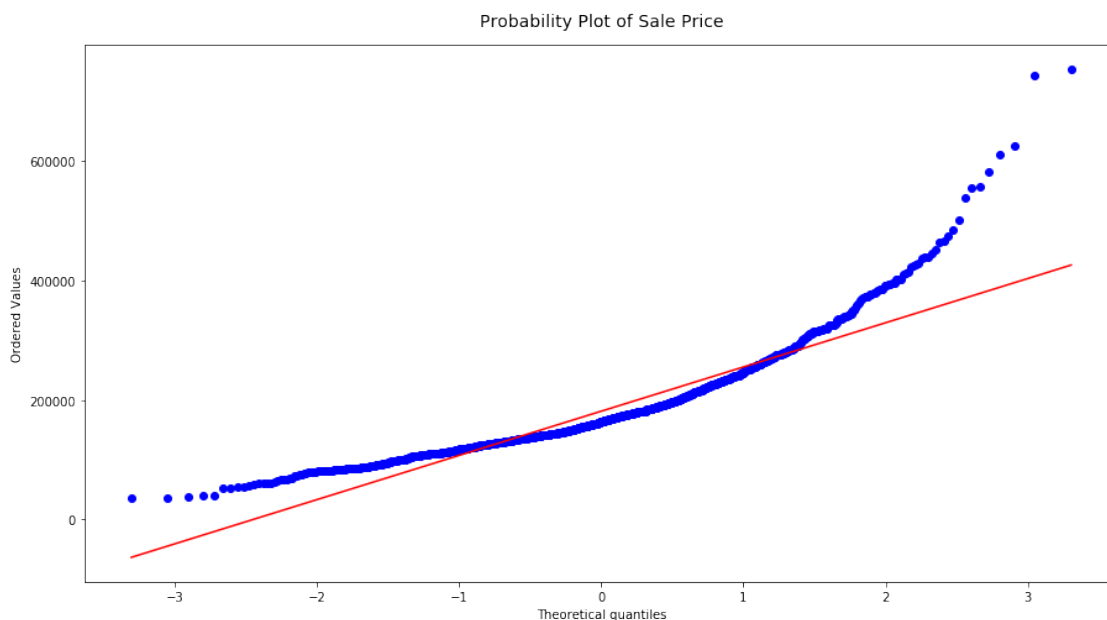
Additionally, I will create a probability plot of sale price to see how the data compares to its best fit line.

```
[13]: plt.figure(figsize = (15,8))

# Plotting quantiles of probability of SalePrice
stats.probplot(train_df['SalePrice'], plot = plt)

plt.title("Probability Plot of Sale Price", fontsize = 14, y=1.02)

plt.show;
```



As we can see in the above graph, SalePrice does not follow the diagonal line, further demonstrating that the data is not normally distributed. One may ask about my obsession with comparing the data to a normal distribution? It is because the simplicity of normally distributed data makes it much easier to build more accurate models. Skewed data can violate inbuilt model assumptions and distort which features are more important in model building. For example, in our housing data set, we may train the model on a much larger number of expensive homes, so when making predictions on moderately priced homes, our model will be less accurate.

5.0.2 Sale Price log-transformation

To transform the distribution of SalePrice into a normal distribution, I will use a log-transformation. This is a simple transformation using the log function from numpy

```
[14]: plt.figure(figsize = (15,8))

#Applying the log function from numpy to all values in the SalePrice column.
→Storing it back in the DF at SalePrice
train_df['SalePrice'] = np.log(train_df.SalePrice)

stats.probplot(train_df['SalePrice'], plot = plt)

plt.title("Log Transformed Probability Plot of Sale Price", fontsize = 14, y=1.
→02)

plt.show;
```



Clearly, the data follows the straight line much better indicating SalePrice has been normalised. We can verify this by calculating the skewness and kurtosis again.

```
[15]: skew_SalePrice = skew(train_df['SalePrice'])
print('Skewness of sale price is:', skew_SalePrice)

kurt_SalePrice = kurtosis(train_df['SalePrice'])
print('Kurtosis of sale price is:', kurt_SalePrice)
```

Skewness of sale price is: 0.1212103673013655

Kurtosis of sale price is: 0.8026555069117713

Therefore, SalePrice is much closer to a normal distribution and is now more useful for when we build models.

Finally, let me replot the distribution of the SalePrice to see its newly normally distributed self in all its glory. In this case, I have compared it to a perfect standard normal distribution of the data

```
[16]: plt.figure(figsize=(15,8))


plt.title("Sale Price Normally Distributed", fontsize = 14)

plt.ylabel("Frequency", fontsize = 14)

plt.xlabel("Sale Price", fontsize = 14)

sns.distplot(train_df.SalePrice, axlabel = False, fit=norm)

plt.show;
```



5.0.3 Missing Values

Now I am going to check which columns have lots of missing values, so I can remove those columns from the dataframe. Columns with lots of missing values shouldn't affect SalePrice much.

```
[17]: #Printing percentage of values missing in each column
missing_lots = train_df.isna().mean() * 100
```


```
#Dropping features with percentage missing values < 6% as these still hold lots
→of information
missing_lots_reduced = missing_lots.drop(missing_lots[missing_lots < 6].index).
→sort_values()
print(missing_lots_reduced)
```

```
LotFrontage    17.739726
FireplaceQu    47.260274
Fence          80.753425
Alley          93.767123
MiscFeature    96.301370
PoolQC         99.520548
dtype: float64
```

```
[18]: # Plotting percentage of missing values
plt.figure(figsize = (15,10))

sns.barplot(x = missing_lots_reduced.index, y = missing_lots_reduced, color =
→'steelblue')

plt.show;
```



As we can see in the plot, the features Alley, PoolQC, Fence and MiscFeatures are missing over

80% of their values. This does not necessarily mean that this data was not collected, but most houses just do not have these features. Because of this, these features will not contribute much to our prediction of SalePrice, hence we can drop them from the Dataframe.

I have decided to drop any features missing more than 15% of their values, as we can conclude that there is not enough information left in that feature to build an accurate model. Furthermore, by looking at the individual features which are missing lots of values, I can assume that they are not very important factors to consider when buying a house, therefore I will drop the above features.

```
[19]: # Dropping features which are missing over 15% of their values
train_df = train_df.drop(columns = missing_lots_reduced.index)
train_df.head()
```

```
[19]:
```

	Id	MSSubClass	MSZoning	LotArea	Street	LotShape	LandContour	Utilities	\
0	1	60	RL	8450	Pave	Reg	Lvl	AllPub	
1	2	20	RL	9600	Pave	Reg	Lvl	AllPub	
2	3	60	RL	11250	Pave	IR1	Lvl	AllPub	
3	4	70	RL	9550	Pave	IR1	Lvl	AllPub	
4	5	60	RL	14260	Pave	IR1	Lvl	AllPub	

	LotConfig	LandSlope	...	EnclosedPorch	3SsnPorch	ScreenPorch	PoolArea	\
0	Inside	Gtl	...	0	0	0	0	
1	FR2	Gtl	...	0	0	0	0	
2	Inside	Gtl	...	0	0	0	0	
3	Corner	Gtl	...	272	0	0	0	
4	FR2	Gtl	...	0	0	0	0	

	MiscVal	MoSold	YrSold	SaleType	SaleCondition	SalePrice
0	0	2	2008	WD	Normal	12.247694
1	0	5	2007	WD	Normal	12.109011
2	0	9	2008	WD	Normal	12.317167
3	0	2	2006	WD	Abnorml	11.849398
4	0	12	2008	WD	Normal	12.429216

[5 rows x 75 columns]

Now I will deal with features with only a few missing values.

```
[20]: missing_small = train_df.isna().mean() * 100

missing_small_reduced = missing_small.drop(missing_small[missing_small == 0].
→index).sort_values()
print(missing_small_reduced)
```

Electrical	0.068493
MasVnrType	0.547945
MasVnrArea	0.547945
BsmtQual	2.534247
BsmtCond	2.534247


```

BsmtFinType1    2.534247
BsmtExposure    2.602740
BsmtFinType2    2.602740
GarageType      5.547945
GarageYrBltd    5.547945
GarageFinish    5.547945
GarageQual      5.547945
GarageCond      5.547945
dtype: float64

```

```

[21]: #Counting the number of missing values for each feature
missing_small_count = train_df.isnull().sum()
missing_small_count_reduced = missing_small_count.
    ↳ drop(missing_small_count[missing_small_count == 0].index).sort_values()
print(missing_small_count_reduced)

```

```

Electrical      1
MasVnrType      8
MasVnrArea      8
BsmtQual       37
BsmtCond       37
BsmtFinType1   37
BsmtExposure   38
BsmtFinType2   38
GarageType     81
GarageYrBltd   81
GarageFinish   81
GarageQual     81
GarageCond     81
dtype: int64

```

Electrical: this feature only has one missing value, so we will just remove that row from the dataframe.

MasVnrType, MasVnrArea, BsmtQual, BsmtCond, BsmtFinType1, BsmtExposure, and BsmtFinType2: all of these features do not have many missing values so we will remove them from the dataframe. Many of these rows overlap so we would only be removing a maximum of 46 rows which only represents 2.6% of the dataframe. So the effect on our model will be minimal.

GarageType: this feature holds useful information about the garage location and if there is no garage so I will impute these 81 values later on.

GarageYrBltd, GarageFinish, GarageQual and GarageCond: these features are all closely related. Hence I will keep GarageQual and remove the rest of the features. The 81 missing values will be imputed later on.

```

[22]: #Removing various rows and features from dataframe
train_df = train_df.dropna(subset = ['Electrical', 'MasVnrType', 'MasVnrArea',
    ↳ 'BsmtQual', 'BsmtCond',

```

```

        'BsmtFinType1', 'BsmtExposure',
        →'BsmtFinType2'])
train_df = train_df.drop(columns = ['GarageYrBlt', 'GarageFinish', 'GarageCond'])

```

```
[23]: train_df.head()
```

```

[23]:   Id  MSSubClass MSZoning  LotArea  Street  LotShape  LandContour  Utilities  \
0    1           60       RL    8450   Pave       Reg         Lvl     AllPub
1    2           20       RL    9600   Pave       Reg         Lvl     AllPub
2    3           60       RL   11250   Pave       IR1         Lvl     AllPub
3    4           70       RL    9550   Pave       IR1         Lvl     AllPub
4    5           60       RL   14260   Pave       IR1         Lvl     AllPub

      LotConfig  LandSlope  ...  EnclosedPorch  3SsnPorch  ScreenPorch  PoolArea  \
0     Inside      Gtl  ...             0             0             0           0
1      FR2      Gtl  ...             0             0             0           0
2     Inside      Gtl  ...             0             0             0           0
3     Corner      Gtl  ...            272             0             0           0
4      FR2      Gtl  ...             0             0             0           0

      MiscVal  MoSold  YrSold  SaleType  SaleCondition  SalePrice
0           0        2    2008         WD         Normal  12.247694
1           0        5    2007         WD         Normal  12.109011
2           0        9    2008         WD         Normal  12.317167
3           0        2    2006         WD        Abnorml  11.849398
4           0       12    2008         WD         Normal  12.429216

```

[5 rows x 72 columns]

5.0.4 Imputation

When I was removing missing values, I was still left with missing values in GarageType and GarageQual. I will now impute these values. Imputation is the process by which you replace a missing value with some other value e.g. mean or in a categorical features case, the most popular type of feature. First, lets get some more information about GarageType and GarageQual.

```

[24]: # Counting the number of instances in a feature
print('No. of instances of values in GarageType: \n', train_df['GarageType'].
      →value_counts(dropna = False))
print('\n No. of instances of values in GarageQual: \n', train_df['GarageQual'].
      →value_counts(dropna = False))

```

```

No. of instances of values in GarageType:
Attchd      852
Detached    369
BuiltIn      85
NaN          74

```

```
Basement      19
CarPort       7
2Types        6
Name: GarageType, dtype: int64
```

```
No. of instances of values in GarageQual:
TA          1270
NaN         74
Fa          48
Gd          14
Ex           3
Po           3
Name: GarageQual, dtype: int64
```

It seems most houses have garages. Therefore for GarageType, I will impute the missing values with Attchd and for GarageQual I will impute the missing values with TA.

```
[25]: #Filling missing values in GarageType with Attchd
train_df['GarageType'].fillna(value = 'Attchd', inplace = True)
```

```
[26]: #Filling missing values in GarageQual with TA
train_df['GarageQual'].fillna(value = 'TA', inplace = True)
```

```
[27]: #Checking that I have dealt with all of the missing values
missing_check = train_df.isna().mean() * 100
missing_check = missing_check.drop(missing_check[missing_check == 0].index).
    ↳sort_values()
print(missing_check)
```

```
Series([], dtype: float64)
```

Clearly, there are no longer any missing values.

5.0.5 Dealing with Outliers

Outliers are pieces of data which do not match the overall distribution of its dataset. These datapoints can arise due to mistakes or variances in the data, and if they detract from the overall information of the dataset then we need to identify and remove them. We can tell if a particular datapoint is an outlier if it lies outside the range of a histogram of all the datapoints of that particular feature. This process is known as univariate analysis.

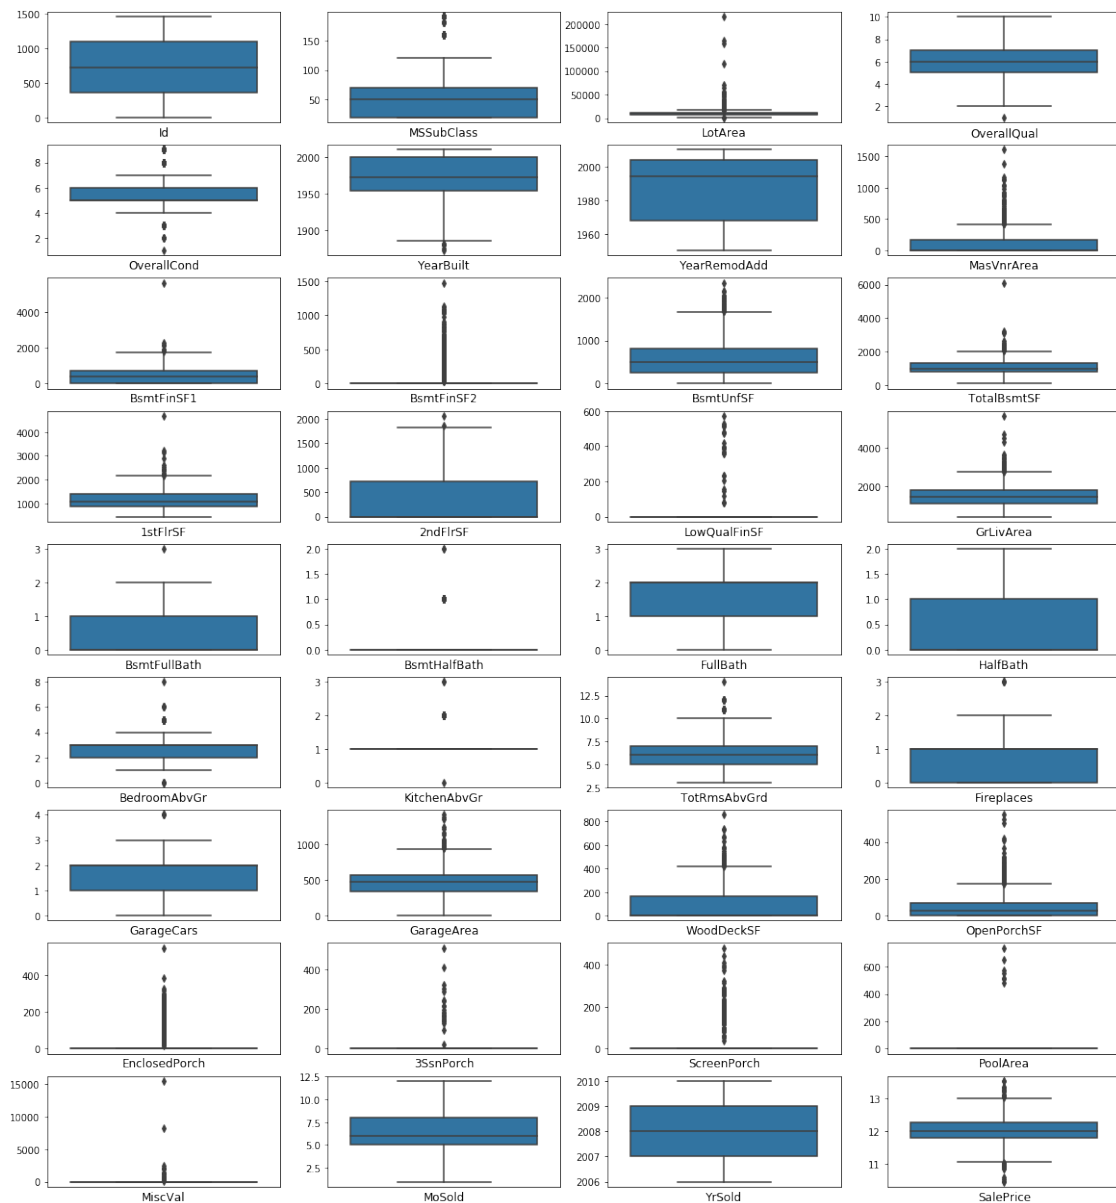
```
[28]: plt.figure(figsize=(20,25))

numerical_cols = train_df.select_dtypes(exclude = 'object')

#This 'for' loop allows me to loop over numerical_cols using enumerate, whilst
    ↳keeping track of the index in i
for i, col in enumerate(numerical_cols.columns):
```

```
# axes allows you to plot multiple axes and do things to the individual axes
→(i.e. label them)
axes = plt.subplot(10,4,i+1)
sns.boxplot(data = train_df[col])
axes.set_xlabel(col, fontsize = 12)
axes.tick_params(axis='x', which='both', bottom=False, top=False,
→labelbottom=False)
```

```
plt.tight_layout
plt.show;
```



As we can see in the graphs above, many of the features display outliers. I am going to remove some of them from the data set to build a more accurate model. However, I will not remove all of the outliers. Some features only have one or two outliers so they have a very minimal impact on the dataset, so it is safe to leave them in. Some outliers are still quite close to the other values in that feature, so do not need to be removed.

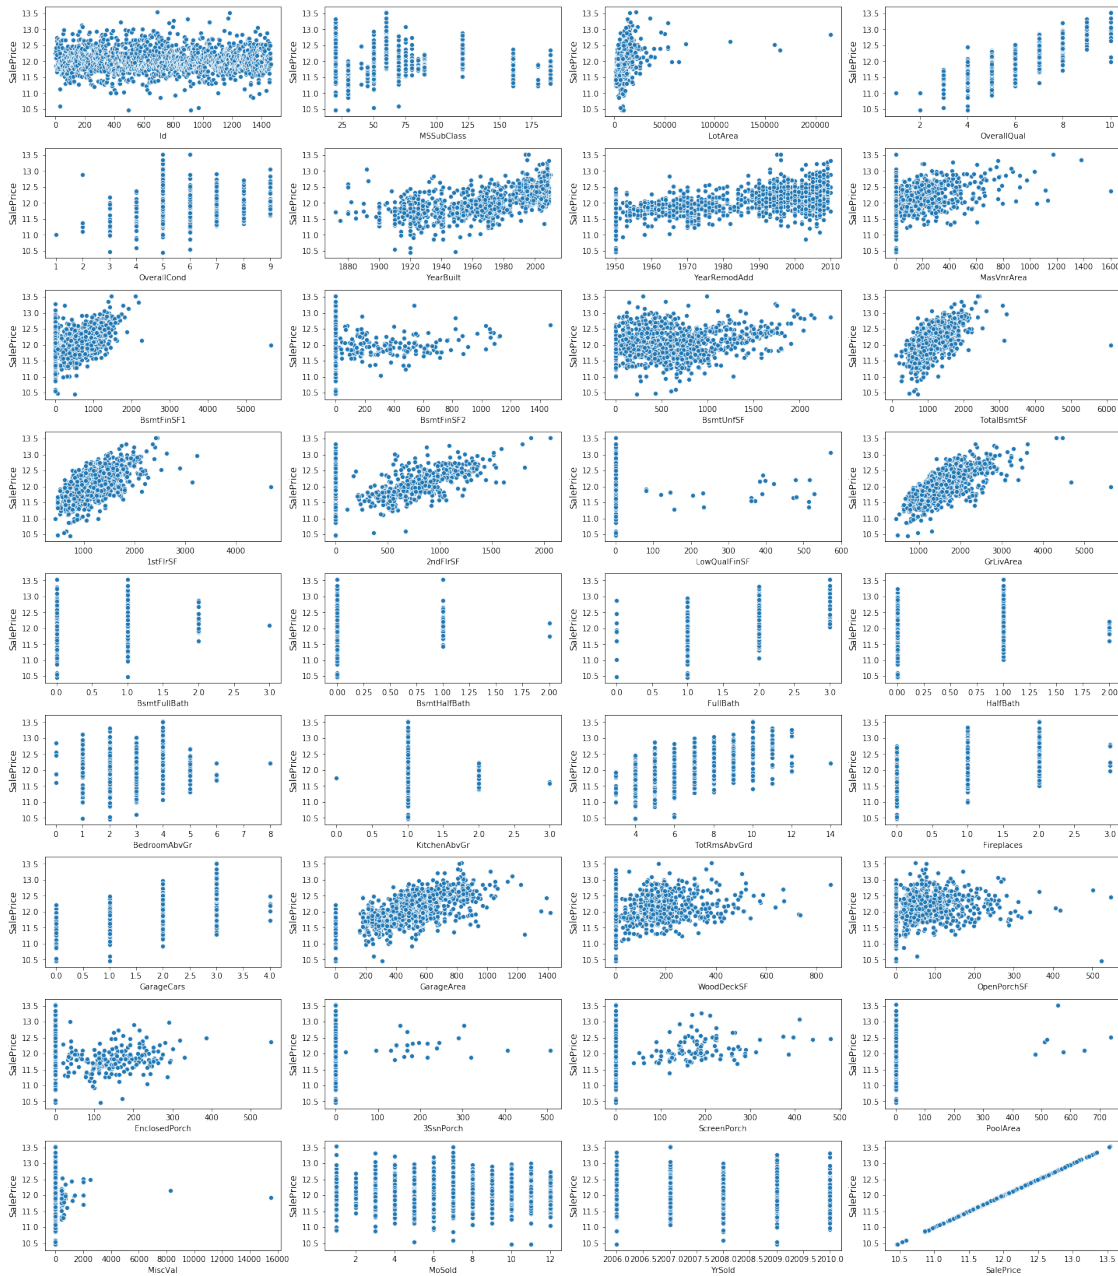
For some of the histograms, it seems they do not have enough data to produce a histogram and demonstrate many outliers (e.g. EnclosedPorch, 3SsnPorch, ScreenPorch, PoolArea, and MiscVal etc). These points are misleading, as they do not really represent outliers, but rather most houses do not have these features hence there not being enough data to produce a histogram. Therefore, these datapoints are not outliers and can remain as part of our dataset.

To choose which outliers from a particular feature which need to be removed, I still need some more information. As the business question is related to SalePrice, I am interested in how these features affect SalePrice. Therefore, I am going to conduct some bivariate analysis and plot scatter plots of the numerical features against SalePrice, to visually see how SalePrice is dependent on each feature.

```
[29]: plt.figure(figsize=(20,25))

for i, col in enumerate(numerical_cols.columns):
    axes = plt.subplot(10,4,i+1)
    sns.scatterplot(x = train_df[col], y = train_df['SalePrice'])
    axes.set_ylabel('SalePrice', fontsize = 12)

plt.tight_layout(pad = 1.0)
plt.show;
```



Based on looking at the number of outliers from the univariate analysis and checking how points vary with SalePrice using bivariate analysis, I have concluded that only the following features need outliers to be removed:

- GrLivArea
- SalePrice

To decide which outliers to remove and to keep in these two features, I will use the 1.5 Interquartile Range (IQR) rule, where any points below (First Quartile - 1.5IQR) and above (Third Quartile + 1.5IQR) will be removed.

```
[30]: #Calculating IQR for GrLivArea and SalePrice
q1_GrLivArea, q3_GrLivArea = np.percentile(train_df['GrLivArea'], [25,75])
iqr_GrLivArea = q3_GrLivArea - q1_GrLivArea

q1_SalePrice, q3_SalePrice = np.percentile(train_df['SalePrice'], [25,75])
iqr_SalePrice = q3_SalePrice - q1_SalePrice
```

```
[31]: #Calculating lower and upper bound thresholds for outliers
lower_bound_GrLivArea = q1_GrLivArea - (1.5 * iqr_GrLivArea)
upper_bound_GrLivArea = q3_GrLivArea + (1.5 * iqr_GrLivArea)

lower_bound_SalePrice = q1_SalePrice - (1.5 * iqr_SalePrice)
upper_bound_SalePrice = q3_SalePrice + (1.5 * iqr_SalePrice)
```

```
[32]: # Creating booleans to find outliers
outliers = train_df[(train_df['GrLivArea'] < lower_bound_GrLivArea) |
    →(train_df['GrLivArea'] > upper_bound_GrLivArea)
    | (train_df['SalePrice'] < lower_bound_SalePrice)
    | (train_df['SalePrice'] > upper_bound_SalePrice)].index

train_df = train_df.drop(outliers)
```

```
[33]: # Checking outlier rows have been dropped
train_df.shape
```

```
[33]: (1362, 72)
```

5.0.6 Encoding Categorical Variables

Most Machine Learning models cannot pass categorical variables into the predictions, therefore I am going to covert them into a usable format. But first, I need to change the type of data represented in some features as they are misleading.

```
[34]: # Changing these features back into strings as currently the dataframe sees them
    →as integers
train_df['MSSubClass'] = str(train_df['MSSubClass'])
train_df['OverallCond'] = str(train_df['OverallCond'])
train_df['OverallQual'] = str(train_df['OverallQual'])
train_df['YrSold'] = str(train_df['YrSold'])
train_df['MoSold'] = str(train_df['MoSold'])
```

```
[35]: train_df.select_dtypes(exclude = 'object').columns # will need later
```

```
[35]: Index(['Id', 'LotArea', 'YearBuilt', 'YearRemodAdd', 'MasVnrArea',
        'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', '1stFlrSF',
        '2ndFlrSF', 'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath',
        'FullBath', 'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'TotRmsAbvGrd',
```

```

'Fireplaces', 'GarageCars', 'GarageArea', 'WoodDeckSF', 'OpenPorchSF',
'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'MiscVal',
'SalePrice'],
dtype='object')

```

I will use a strategy called one-hot encoding which gives a value of 1 for each type which occurs in a particular categorical feature and gives a value of 0 if that type is not present. This is a simple way to encode all of the categorical features and should prove useful

```

[36]: train_df = pd.get_dummies(train_df)
      train_df.head()

```

```

[36]:   Id  LotArea  YearBuilt  YearRemodAdd  MasVnrArea  BsmtFinSF1  BsmtFinSF2  \
0    1    8450    2003    2003    196.0    706    0
1    2    9600    1976    1976     0.0    978    0
2    3   11250    2001    2002    162.0    486    0
3    4    9550    1915    1970     0.0    216    0
4    5   14260    2000    2000    350.0    655    0

      BsmtUnfSF  TotalBsmtSF  1stFlrSF  ...  SaleType_ConLw  SaleType_New  \
0         150         856      856  ...          0          0
1         284        1262     1262  ...          0          0
2         434         920      920  ...          0          0
3         540         756      961  ...          0          0
4         490        1145     1145  ...          0          0

      SaleType_0th  SaleType_WD  SaleCondition_Abnorml  SaleCondition_AdjLand  \
0              0              1              0              0
1              0              1              0              0
2              0              1              0              0
3              0              1              1              0
4              0              1              0              0

      SaleCondition_Alloca  SaleCondition_Family  SaleCondition_Normal  \
0              0              0              1
1              0              0              1
2              0              0              1
3              0              0              0
4              0              0              1

      SaleCondition_Partial
0              0
1              0
2              0
3              0
4              0

```


[5 rows x 255 columns]

5.0.7 Checking skewness and kurtosis of rest of dataset

Only the numerical features demonstrate skewness and kurtosis, so we will select these columns first. And then we will check their skewness and kurtosis to decide whether we will log transform the feature. NB: we will not do this for categorical features which have been encoded. I will only use the numerical features which are not discrete (e.g. YearBuilt etc.)

```
[37]: numerical_columns = ['LotArea', 'MasVnrArea',  
                          'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', '1stFlrSF',  
                          '2ndFlrSF', 'LowQualFinSF', 'GrLivArea',  
                          'BedroomAbvGr', 'GarageArea', 'WoodDeckSF', 'OpenPorchSF',  
                          'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'MiscVal']  
  
for i, col in enumerate(numerical_columns):  
    skewness = train_df[col].skew()  
    kurt = train_df[col].kurtosis()  
    print('Skewness of', col, 'is', skewness)  
    print('Kurtosis of', col, 'is', kurt, '\n')
```

Skewness of LotArea is 12.728694101002754

Kurtosis of LotArea is 213.67938372752775

Skewness of MasVnrArea is 2.590881997294897

Kurtosis of MasVnrArea is 10.073190763893653

Skewness of BsmtFinSF1 is 0.6687200389155128

Kurtosis of BsmtFinSF1 is -0.3804752689892852

Skewness of BsmtFinSF2 is 4.222610786234709

Kurtosis of BsmtFinSF2 is 19.68631550787231

Skewness of BsmtUnfSF is 0.9304248401095043

Kurtosis of BsmtUnfSF is 0.5284944356251353

Skewness of TotalBsmtSF is 0.69746052606669

Kurtosis of TotalBsmtSF is 0.26837006104123384

Skewness of 1stFlrSF is 0.697982873735981

Kurtosis of 1stFlrSF is 0.10848807187995524

Skewness of 2ndFlrSF is 0.7119746132221992

Kurtosis of 2ndFlrSF is -0.9668470217148708

Skewness of LowQualFinSF is 10.161640363702475

Kurtosis of LowQualFinSF is 107.23029942840147

Skewness of GrLivArea is 0.4829127905098185
Kurtosis of GrLivArea is -0.1948566522712274

Skewness of BedroomAbvGr is 0.008206513816296715
Kurtosis of BedroomAbvGr is 1.6499527281015074

Skewness of GarageArea is 0.17071879537395862
Kurtosis of GarageArea is 0.90699530930583

Skewness of WoodDeckSF is 1.5377435286917205
Kurtosis of WoodDeckSF is 3.113487147788757

Skewness of OpenPorchSF is 1.874472931805437
Kurtosis of OpenPorchSF is 4.1642698234699544

Skewness of EnclosedPorch is 2.9089863253866723
Kurtosis of EnclosedPorch is 7.913576535857889

Skewness of 3SsnPorch is 10.189015567275886
Kurtosis of 3SsnPorch is 120.48183060568954

Skewness of ScreenPorch is 3.8333486778336594
Kurtosis of ScreenPorch is 14.847027314873504

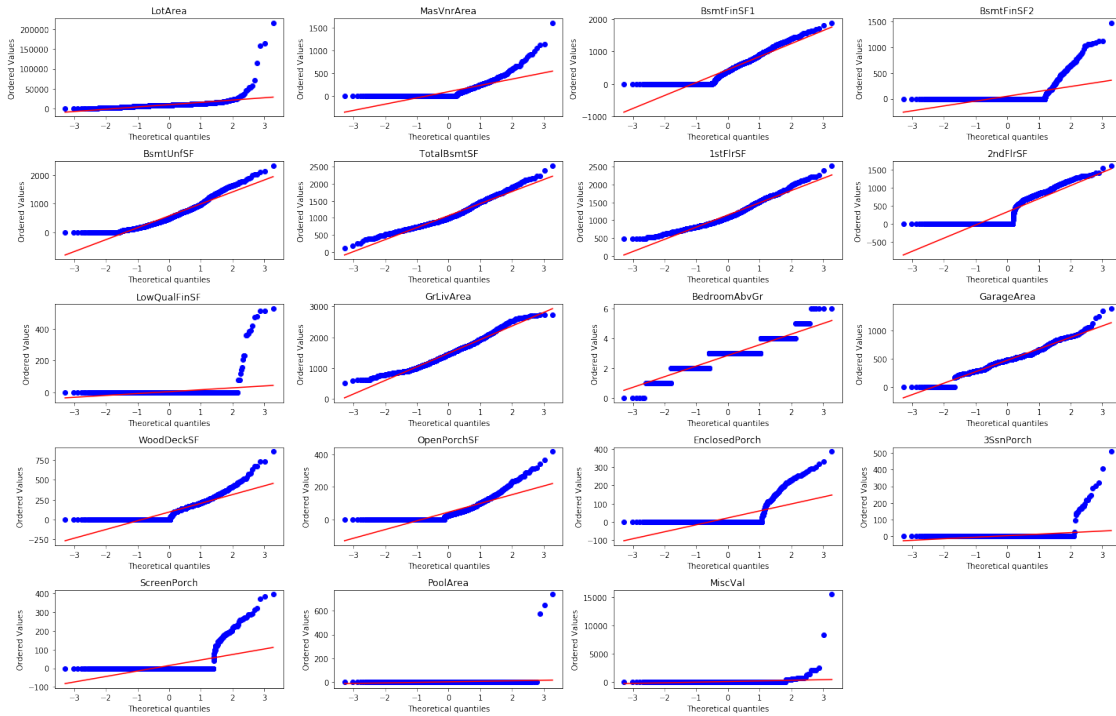
Skewness of PoolArea is 21.586982953398458
Kurtosis of PoolArea is 469.29174240478625

Skewness of MiscVal is 25.334782804784542
Kurtosis of MiscVal is 726.9251286699115

```
[38]: plt.figure(figsize = (20,25))

for i, col in enumerate(numerical_columns):
    axes = plt.subplot(10,4,i+1)
    stats.probplot(train_df[col], plot = plt)
    plt.title(col)

plt.tight_layout(pad = 1.0)
```



Ok, now we have the skewness, kurtosis and probability plots of the numerical features, let's decide which features need to be log-transformed. As ideal values of skewness and kurtosis are both 0, some features have skewness and kurtosis quite close to this value. Hence I will define a threshold of skewness of above 0.75 for features to be log transformed. As features who have low skewness also have a low kurtosis, we do not need to set a threshold for kurtosis.

```
[39]: numerical_columns = ['LotArea', 'MasVnrArea',
    'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', '1stFlrSF',
    '2ndFlrSF', 'LowQualFinSF', 'GrLivArea',
    'BedroomAbvGr', 'GarageArea', 'WoodDeckSF', 'OpenPorchSF',
    'EnclosedPorch', '3SsnPorch', 'ScreenPorch']
```

```
def high_skew(dataframe, colls):
    """
    This function will only keep features which have high skewness

    Input: Dataframe

    Output: Array with features which have high skew

    """

    #Creating empty array to store high skewed features
```

```

high_skewed_features = []

for i, col in enumerate(colls):

    skewness_df = dataframe[col].skew()

    #Condition for which features to keep
    if skewness_df > 1:

        #Adding the feature name to array
        high_skewed_features.append(col)

return high_skewed_features

```

```

[40]: #Calling function on train_df
high_skewed_features_train_df = high_skew(train_df, numerical_columns)
print(high_skewed_features_train_df)

```

```

['LotArea', 'MasVnrArea', 'BsmtFinSF2', 'LowQualFinSF', 'WoodDeckSF',
'OpenPorchSF', 'EnclosedPorch', '3SsnPorch', 'ScreenPorch']

```

```

[41]: import scipy

#Performing a boxcox transformation to manage skewness
high_skewed_features_train_df_box = scipy.special.
    ↪boxcox1p(train_df[high_skewed_features_train_df], 0.5)

```

Unlike before when I used a log transformation, now my features have negative values hence a log transformation does not work. Therefore, I use a boxcox transformation. For more information about how boxcox transformations: <https://docs.scipy.org/doc/scipy/reference/generated/scipy.special.boxcox1p.html>

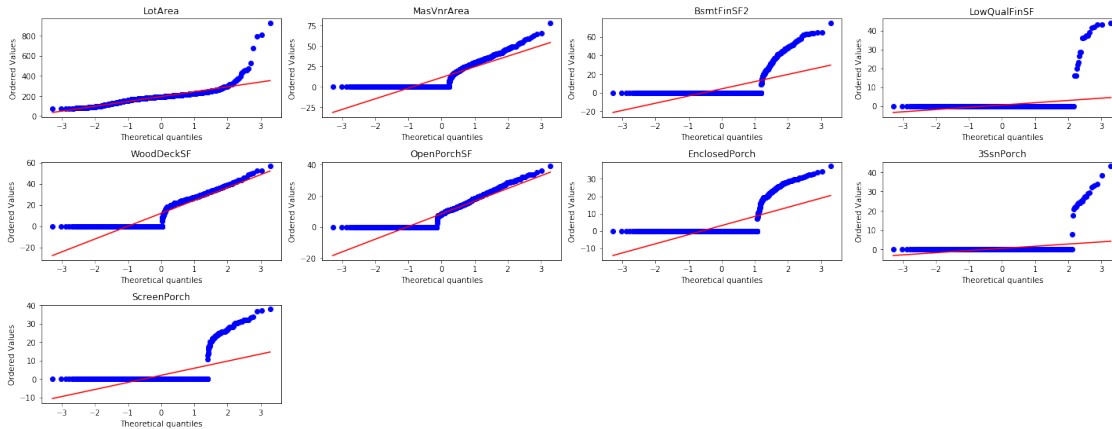
```

[42]: plt.figure(figsize = (20,25))

#Checking new lines of best fit
for i, col in enumerate(high_skewed_features_train_df):
    axes = plt.subplot(10,4,i+1)
    stats.probplot(high_skewed_features_train_df_box[col], plot = plt)
    plt.title(col)

plt.tight_layout(pad = 1.0)

```



```
[43]: #Checking skewness is reduced
for i, col in enumerate(high_skewed_features_train_df):
    skewness = high_skewed_features_train_df_box[col].skew()
    print('Skewness of', col, 'is', skewness)
```

```
Skewness of LotArea is 4.368044693263819
Skewness of MasVnrArea is 1.0535272749685074
Skewness of BsmtFinSF2 is 3.0588149311390076
Skewness of LowQualFinSF is 9.017708889454473
Skewness of WoodDeckSF is 0.5064592144802054
Skewness of OpenPorchSF is 0.6400877188563249
Skewness of EnclosedPorch is 2.381578795644092
Skewness of 3SsnPorch is 8.270539021633104
Skewness of ScreenPorch is 3.30897987124542
```

As we can see the skewness has been reduced for the high skew features. Although some still have high skews (skewness > 5), this could be because of a number of reasons, such as less houses having this feature or the houses which have these features are very different to the rest of the dataset. Therefore, as only a few features demonstrate this high skewness, we can just ignore it as it won't have a large effect on our model.

```
[44]: #Dropping columns with skewed values and replacing them with transformed values

train_df[high_skewed_features_train_df] = high_skewed_features_train_df_box
```

5.0.8 And beyond...

Now my training dataset is finally ready for modelling. Let me show my dataframe one last time.

```
[45]: train_df
```

```
[45]:      Id    LotArea  YearBuilt  YearRemodAdd  MasVnrArea  BsmtFinSF1  \
0      1  181.858641      2003           2003    26.071338      706
```

1	2	193.969385	1976	1976	0.000000	978
2	3	210.141462	2001	2002	23.534291	486
3	4	193.458435	1915	1970	0.000000	216
4	5	236.838858	2000	2000	35.469988	655
...
1455	1456	175.966289	1999	2000	0.000000	0
1456	1457	227.573518	1978	1988	19.908902	790
1457	1458	188.189379	1941	2006	0.000000	275
1458	1459	195.159834	1950	1996	0.000000	49
1459	1460	197.379036	1965	1965	0.000000	830

	BsmtFinSF2	BsmtUnfSF	TotalBsmtSF	1stFlrSF	...	SaleType_ConLw	\
0	0.000000	150	856	856	...	0	
1	0.000000	284	1262	1262	...	0	
2	0.000000	434	920	920	...	0	
3	0.000000	540	756	961	...	0	
4	0.000000	490	1145	1145	...	0	
...	
1455	0.000000	953	953	953	...	0	
1456	23.612497	589	1542	2073	...	0	
1457	0.000000	877	1152	1188	...	0	
1458	62.187226	0	1078	1078	...	0	
1459	32.117444	136	1256	1256	...	0	

	SaleType_New	SaleType_Oth	SaleType_WD	SaleCondition_Abnorml	\
0	0	0	1	0	
1	0	0	1	0	
2	0	0	1	0	
3	0	0	1	1	
4	0	0	1	0	
...	
1455	0	0	1	0	
1456	0	0	1	0	
1457	0	0	1	0	
1458	0	0	1	0	
1459	0	0	1	0	

	SaleCondition_AdjLand	SaleCondition_Alloca	SaleCondition_Family	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	
...	
1455	0	0	0	
1456	0	0	0	
1457	0	0	0	

1458	0	0	0
1459	0	0	0

	SaleCondition_Normal	SaleCondition_Partial
0	1	0
1	1	0
2	1	0
3	0	0
4	1	0
...
1455	1	0
1456	1	0
1457	1	0
1458	1	0
1459	1	0

[1362 rows x 255 columns]

6 Cleaning Test Data

This section will use many of the techniques and decisions I made when processing the training dataset.

```
[46]: test_df
```

```
[46]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	\
0	1461	20	RH	80.0	11622	Pave	NaN	Reg	
1	1462	20	RL	81.0	14267	Pave	NaN	IR1	
2	1463	60	RL	74.0	13830	Pave	NaN	IR1	
3	1464	60	RL	78.0	9978	Pave	NaN	IR1	
4	1465	120	RL	43.0	5005	Pave	NaN	IR1	
...	
1454	2915	160	RM	21.0	1936	Pave	NaN	Reg	
1455	2916	160	RM	21.0	1894	Pave	NaN	Reg	
1456	2917	20	RL	160.0	20000	Pave	NaN	Reg	
1457	2918	85	RL	62.0	10441	Pave	NaN	Reg	
1458	2919	60	RL	74.0	9627	Pave	NaN	Reg	

	LandContour	Utilities	...	ScreenPorch	PoolArea	PoolQC	Fence	\
0	Lvl	AllPub	...	120	0	NaN	MnPrv	
1	Lvl	AllPub	...	0	0	NaN	NaN	
2	Lvl	AllPub	...	0	0	NaN	MnPrv	
3	Lvl	AllPub	...	0	0	NaN	NaN	
4	HLS	AllPub	...	144	0	NaN	NaN	
...	
1454	Lvl	AllPub	...	0	0	NaN	NaN	

1455	Lvl	AllPub	...	0	0	NaN	NaN
1456	Lvl	AllPub	...	0	0	NaN	NaN
1457	Lvl	AllPub	...	0	0	NaN	MnPrv
1458	Lvl	AllPub	...	0	0	NaN	NaN

	MiscFeature	MiscVal	MoSold	YrSold	SaleType	SaleCondition
0	NaN	0	6	2010	WD	Normal
1	Gar2	12500	6	2010	WD	Normal
2	NaN	0	3	2010	WD	Normal
3	NaN	0	6	2010	WD	Normal
4	NaN	0	1	2010	WD	Normal
...
1454	NaN	0	6	2006	WD	Normal
1455	NaN	0	4	2006	WD	Abnorml
1456	NaN	0	9	2006	WD	Abnorml
1457	Shed	700	7	2006	WD	Normal
1458	NaN	0	11	2006	WD	Normal

[1459 rows x 80 columns]

Lets convert the values which are currently not in their correct format

```
[47]: test_df['MSSubClass'] = str(test_df['MSSubClass'])
test_df['OverallCond'] = str(test_df['OverallCond'])
test_df['OverallQual'] = str(test_df['OverallQual'])
test_df['YrSold'] = str(test_df['YrSold'])
test_df['MoSold'] = str(test_df['MoSold'])
```

6.0.1 Missing Values

```
[48]: test_df.isnull().sum()
```

```
[48]: Id                0
MSSubClass            0
MSZoning              4
LotFrontage         227
LotArea              0
...
MiscVal              0
MoSold              0
YrSold              0
SaleType            1
SaleCondition        0
Length: 80, dtype: int64
```

This is not that useful, lets only retrieve the columns with missing values and what percentage of values are missing.


```
[49]: test_missing = test_df.isna().mean() * 100

test_missing_drop = test_missing.drop(test_missing[test_missing == 0].index).
    ↪sort_values()
print(test_missing_drop)
```

```
TotalBsmtSF      0.068540
GarageArea        0.068540
GarageCars        0.068540
KitchenQual       0.068540
BsmtUnfSF         0.068540
BsmtFinSF2        0.068540
BsmtFinSF1        0.068540
SaleType          0.068540
Exterior1st       0.068540
Exterior2nd       0.068540
Functional        0.137080
Utilities         0.137080
BsmtHalfBath      0.137080
BsmtFullBath      0.137080
MSZoning          0.274160
MasVnrArea        1.028101
MasVnrType        1.096642
BsmtFinType2      2.878684
BsmtFinType1      2.878684
BsmtQual          3.015764
BsmtExposure      3.015764
BsmtCond          3.084304
GarageType        5.209047
GarageFinish      5.346127
GarageQual        5.346127
GarageCond        5.346127
GarageYrBlt       5.346127
LotFrontage       15.558602
FireplaceQu       50.034270
Fence             80.123372
Alley             92.666210
MiscFeature       96.504455
PoolQC           99.794380
dtype: float64
```

As we want to pass the same features from the test set into our model which was fitted on the training set, let's see which columns in the test set we can remove straight away.

```
[50]: train_df[train_df.columns[0:30]].columns
```

```
[50]: Index(['Id', 'LotArea', 'YearBuilt', 'YearRemodAdd', 'MasVnrArea',
          'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', '1stFlrSF',
          '2ndFlrSF', 'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath',
          'FullBath', 'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'TotRmsAbvGrd',
          'Fireplaces', 'GarageCars', 'GarageArea', 'WoodDeckSF', 'OpenPorchSF',
          'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'MiscVal'],
          dtype='object')
```

```
[51]: test_df.select_dtypes(include = np.number).columns
```

```
[51]: Index(['Id', 'LotFrontage', 'LotArea', 'YearBuilt', 'YearRemodAdd',
          'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF',
          '1stFlrSF', '2ndFlrSF', 'LowQualFinSF', 'GrLivArea', 'BsmtFullBath',
          'BsmtHalfBath', 'FullBath', 'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr',
          'TotRmsAbvGrd', 'Fireplaces', 'GarageYrBlt', 'GarageCars', 'GarageArea',
          'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch', '3SsnPorch',
          'ScreenPorch', 'PoolArea', 'MiscVal'],
          dtype='object')
```

Ok, so we can drop 'Id', 'LotFrontage', 'GarageYrBlt'.

```
[52]: test_df = test_df.drop(columns = ['Id', 'LotFrontage', 'GarageYrBlt'])
```

```
[53]: test_missing = test_df.isna().mean() * 100

test_missing_drop = test_missing.drop(test_missing[test_missing == 0].index).
    ↪sort_values()
print(test_missing_drop)
```

TotalBsmtSF	0.068540
GarageArea	0.068540
GarageCars	0.068540
KitchenQual	0.068540
BsmtUnfSF	0.068540
BsmtFinSF2	0.068540
BsmtFinSF1	0.068540
SaleType	0.068540
Exterior1st	0.068540
Exterior2nd	0.068540
Functional	0.137080
Utilities	0.137080
BsmtHalfBath	0.137080
BsmtFullBath	0.137080
MSZoning	0.274160
MasVnrArea	1.028101
MasVnrType	1.096642
BsmtFinType2	2.878684
BsmtFinType1	2.878684

BsmtQual	3.015764
BsmtExposure	3.015764
BsmtCond	3.084304
GarageType	5.209047
GarageFinish	5.346127
GarageQual	5.346127
GarageCond	5.346127
FireplaceQu	50.034270
Fence	80.123372
Alley	92.666210
MiscFeature	96.504455
PoolQC	99.794380

dtype: float64

We can remove 'Fence', 'Alley', 'MiscFeature', 'PoolQC' as they have lots of missing values so do not contribute much to the model.

```
[54]: test_df = test_df.drop(columns = ['Fence', 'Alley', 'MiscFeature', 'PoolQC'],
    ↪ 'FireplaceQu'])
```

```
[55]: test_missing = test_df.isna().mean() * 100

test_missing_drop = test_missing.drop(test_missing[test_missing == 0].index).
    ↪ sort_values()
print(test_missing_drop)
```

BsmtFinSF2	0.068540
GarageArea	0.068540
GarageCars	0.068540
KitchenQual	0.068540
TotalBsmtSF	0.068540
BsmtUnfSF	0.068540
BsmtFinSF1	0.068540
SaleType	0.068540
Exterior1st	0.068540
Exterior2nd	0.068540
BsmtHalfBath	0.137080
Utilities	0.137080
Functional	0.137080
BsmtFullBath	0.137080
MSZoning	0.274160
MasVnrArea	1.028101
MasVnrType	1.096642
BsmtFinType2	2.878684
BsmtFinType1	2.878684
BsmtQual	3.015764
BsmtExposure	3.015764
BsmtCond	3.084304

```
GarageType      5.209047
GarageCond      5.346127
GarageFinish    5.346127
GarageQual      5.346127
dtype: float64
```

Now I am going to check which of these columns I have already removed from the training set so I can remove these columns. Then I will impute the rest of the missing values with 0.

These columns can be removed as they are not in our training set 'BsmtQual', 'GarageCond', 'GarageFinish'.

```
[56]: test_df = test_df.drop(columns = ['BsmtQual', 'GarageCond', 'GarageFinish'])
```

```
[57]: test_missing = test_df.isna().mean() * 100

test_missing_drop = test_missing.drop(test_missing[test_missing == 0].index).
    ↪sort_values()
print(test_missing_drop)
```

```
BsmtFinSF2      0.068540
GarageArea      0.068540
GarageCars      0.068540
KitchenQual     0.068540
TotalBsmtSF     0.068540
BsmtUnfSF       0.068540
BsmtFinSF1      0.068540
SaleType        0.068540
Exterior2nd     0.068540
Exterior1st     0.068540
BsmtFullBath    0.137080
BsmtHalfBath    0.137080
Functional      0.137080
Utilities       0.137080
MSZoning        0.274160
MasVnrArea     1.028101
MasVnrType      1.096642
BsmtFinType1    2.878684
BsmtFinType2    2.878684
BsmtExposure    3.015764
BsmtCond        3.084304
GarageType      5.209047
GarageQual      5.346127
dtype: float64
```

```
[58]: test_df = test_df.fillna(0)
```

```
[59]: test_missing = test_df.isna().mean() * 100

test_missing_drop = test_missing.drop(test_missing[test_missing == 0].index).
    ↪sort_values()
print(test_missing_drop)
```

Series([], dtype: float64)

Now we no longer have any missing values.

```
[60]: test_df
```

```
[60]:
```

				MSSubClass	MSZoning	LotArea	\
0	0	20\n1	20\n2	60\n3	...	RH	11622
1	0	20\n1	20\n2	60\n3	...	RL	14267
2	0	20\n1	20\n2	60\n3	...	RL	13830
3	0	20\n1	20\n2	60\n3	...	RL	9978
4	0	20\n1	20\n2	60\n3	...	RL	5005
...				
1454	0	20\n1	20\n2	60\n3	...	RM	1936
1455	0	20\n1	20\n2	60\n3	...	RM	1894
1456	0	20\n1	20\n2	60\n3	...	RL	20000
1457	0	20\n1	20\n2	60\n3	...	RL	10441
1458	0	20\n1	20\n2	60\n3	...	RL	9627

	Street	LotShape	LandContour	Utilities	LotConfig	LandSlope	Neighborhood	\
0	Pave	Reg	Lvl	AllPub	Inside	Gtl	Names	
1	Pave	IR1	Lvl	AllPub	Corner	Gtl	Names	
2	Pave	IR1	Lvl	AllPub	Inside	Gtl	Gilbert	
3	Pave	IR1	Lvl	AllPub	Inside	Gtl	Gilbert	
4	Pave	IR1	HLS	AllPub	Inside	Gtl	StoneBr	
...	
1454	Pave	Reg	Lvl	AllPub	Inside	Gtl	MeadowV	
1455	Pave	Reg	Lvl	AllPub	Inside	Gtl	MeadowV	
1456	Pave	Reg	Lvl	AllPub	Inside	Gtl	Mitchel	
1457	Pave	Reg	Lvl	AllPub	Inside	Gtl	Mitchel	
1458	Pave	Reg	Lvl	AllPub	Inside	Mod	Mitchel	

	...	OpenPorchSF	EnclosedPorch	3SsnPorch	ScreenPorch	PoolArea	MiscVal	\
0	...	0	0	0	120	0	0	
1	...	36	0	0	0	0	12500	
2	...	34	0	0	0	0	0	
3	...	36	0	0	0	0	0	
4	...	82	0	0	144	0	0	
...	
1454	...	0	0	0	0	0	0	
1455	...	24	0	0	0	0	0	
1456	...	0	0	0	0	0	0	

1457	...	32	0	0	0	0	700
1458	...	48	0	0	0	0	0

					MoSold \
0	0	6\n1	6\n2	3\n3	6...
1	0	6\n1	6\n2	3\n3	6...
2	0	6\n1	6\n2	3\n3	6...
3	0	6\n1	6\n2	3\n3	6...
4	0	6\n1	6\n2	3\n3	6...
...					...
1454	0	6\n1	6\n2	3\n3	6...
1455	0	6\n1	6\n2	3\n3	6...
1456	0	6\n1	6\n2	3\n3	6...
1457	0	6\n1	6\n2	3\n3	6...
1458	0	6\n1	6\n2	3\n3	6...

					YrSold	SaleType	SaleCondition
0	0	2010\n1	2010\n2	2010\n3	...	WD	Normal
1	0	2010\n1	2010\n2	2010\n3	...	WD	Normal
2	0	2010\n1	2010\n2	2010\n3	...	WD	Normal
3	0	2010\n1	2010\n2	2010\n3	...	WD	Normal
4	0	2010\n1	2010\n2	2010\n3	...	WD	Normal
...				
1454	0	2010\n1	2010\n2	2010\n3	...	WD	Normal
1455	0	2010\n1	2010\n2	2010\n3	...	WD	Abnorml
1456	0	2010\n1	2010\n2	2010\n3	...	WD	Abnorml
1457	0	2010\n1	2010\n2	2010\n3	...	WD	Normal
1458	0	2010\n1	2010\n2	2010\n3	...	WD	Normal

[1459 rows x 69 columns]

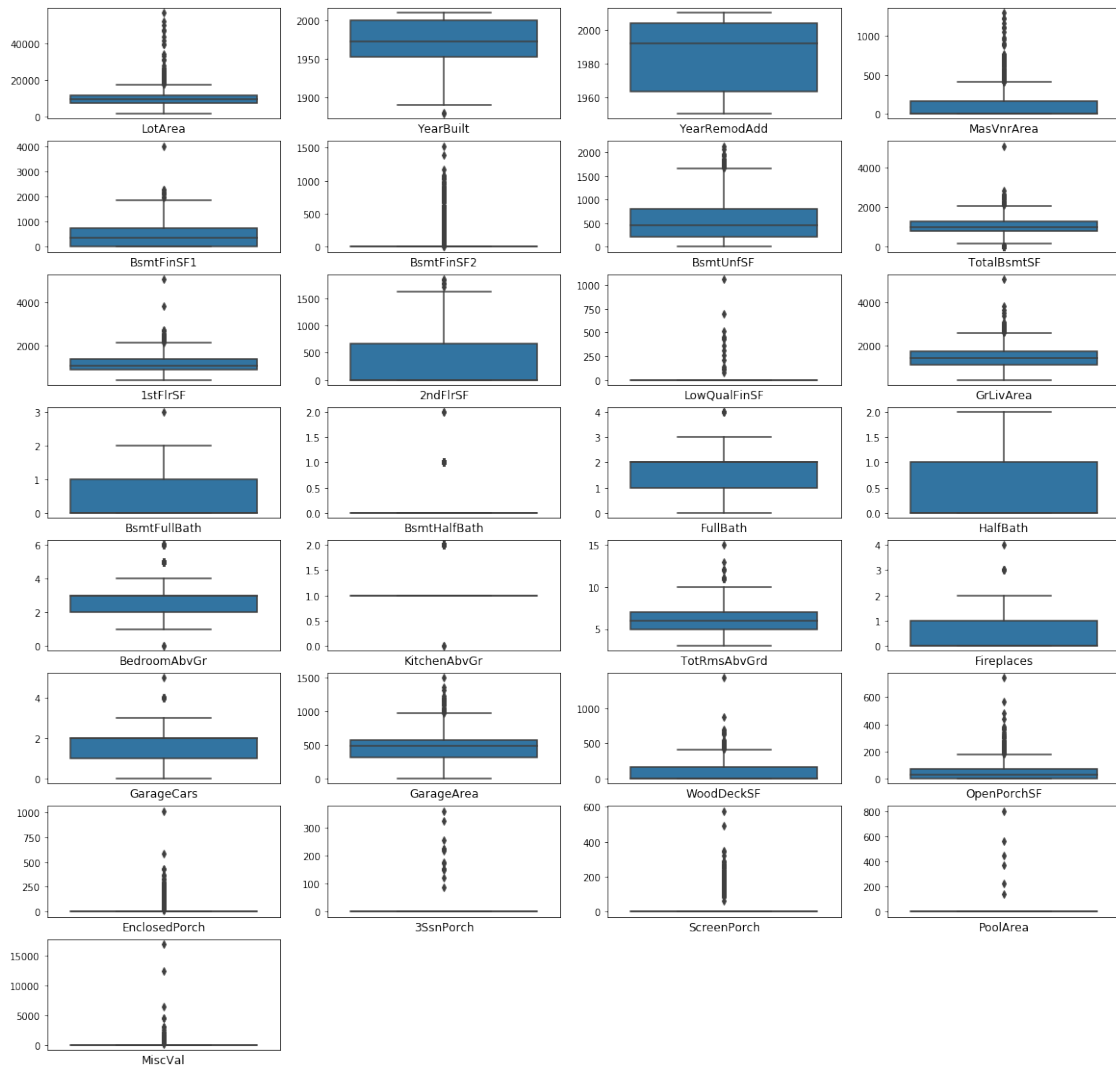
6.0.2 Outliers

```
[61]: plt.figure(figsize=(20,25))

numerical_cols = test_df.select_dtypes(exclude = 'object')

#This 'for' loop allows me to loop over numerical_cols using enumerate, whilst
→keeping track of the index in i
for i, col in enumerate(numerical_cols.columns):
    # axes allows you to plot multiple axes and do things to the individual axes
    →(i.e. label them)
    axes = plt.subplot(10,4,i+1)
    sns.boxplot(data = test_df[col])
    axes.set_xlabel(col, fontsize = 12)
    axes.tick_params(axis='x', which='both', bottom=False, top=False,
    →labelbottom=False)
```

```
plt.tight_layout
plt.show;
```



Clearly, 'LotArea', 'MasVnrArea', 'GrLivArea', 'GarageArea', 'OpenPorchSF' have outliers which can be removed.

```
[62]: def outlier_drop(data_frame, columns_outliers):
    """
    This function will drop all of the outliers in features which are specified
    Input: dataframe and features with outliers that need removing
    Output: updated dataframe with outliers removed
```

```

"""

for i, col in enumerate(columns_outliers):
    q1_col, q3_col = np.percentile(data_frame[col], [25,75])
    iqr_col = q3_col - q1_col

    lowerbound_col = q1_col - (1.5 * iqr_col)
    upperbound_col = q3_col + (1.5 * iqr_col)

    outliers = (data_frame[(data_frame[col] < lowerbound_col)
                           | (data_frame[col] > upperbound_col)].index)

    data_frame = data_frame.drop(outliers)

return data_frame

```

```

[63]: outliers_columns = ['LotArea', 'MasVnrArea', 'GrLivArea', 'GarageArea',
    ↪ 'OpenPorchSF']

#Removing outliers from test dataframe
outlier_drop(test_df, outliers_columns)

```

```

[63]:

```

					MSSubClass	MSZoning	LotArea	\
0	0	20\n1	20\n2	60\n3	...	RH	11622	
1	0	20\n1	20\n2	60\n3	...	RL	14267	
2	0	20\n1	20\n2	60\n3	...	RL	13830	
3	0	20\n1	20\n2	60\n3	...	RL	9978	
4	0	20\n1	20\n2	60\n3	...	RL	5005	
...					
1454	0	20\n1	20\n2	60\n3	...	RM	1936	
1455	0	20\n1	20\n2	60\n3	...	RM	1894	
1456	0	20\n1	20\n2	60\n3	...	RL	20000	
1457	0	20\n1	20\n2	60\n3	...	RL	10441	
1458	0	20\n1	20\n2	60\n3	...	RL	9627	

	Street	LotShape	LandContour	Utilities	LotConfig	LandSlope	Neighborhood	\
0	Pave	Reg	Lvl	AllPub	Inside	Gtl	Names	
1	Pave	IR1	Lvl	AllPub	Corner	Gtl	Names	
2	Pave	IR1	Lvl	AllPub	Inside	Gtl	Gilbert	
3	Pave	IR1	Lvl	AllPub	Inside	Gtl	Gilbert	
4	Pave	IR1	HLS	AllPub	Inside	Gtl	StoneBr	
...	
1454	Pave	Reg	Lvl	AllPub	Inside	Gtl	MeadowV	
1455	Pave	Reg	Lvl	AllPub	Inside	Gtl	MeadowV	
1456	Pave	Reg	Lvl	AllPub	Inside	Gtl	Mitchel	
1457	Pave	Reg	Lvl	AllPub	Inside	Gtl	Mitchel	
1458	Pave	Reg	Lvl	AllPub	Inside	Mod	Mitchel	

	...	OpenPorchSF	EnclosedPorch	3SsnPorch	ScreenPorch	PoolArea	MiscVal	\
0	...	0	0	0	120	0	0	
1	...	36	0	0	0	0	12500	
2	...	34	0	0	0	0	0	
3	...	36	0	0	0	0	0	
4	...	82	0	0	144	0	0	
...	
1454	...	0	0	0	0	0	0	
1455	...	24	0	0	0	0	0	
1456	...	0	0	0	0	0	0	
1457	...	32	0	0	0	0	700	
1458	...	48	0	0	0	0	0	

					MoSold	\
0	0	6\n1	6\n2	3\n3	6...	
1	0	6\n1	6\n2	3\n3	6...	
2	0	6\n1	6\n2	3\n3	6...	
3	0	6\n1	6\n2	3\n3	6...	
4	0	6\n1	6\n2	3\n3	6...	
...					...	
1454	0	6\n1	6\n2	3\n3	6...	
1455	0	6\n1	6\n2	3\n3	6...	
1456	0	6\n1	6\n2	3\n3	6...	
1457	0	6\n1	6\n2	3\n3	6...	
1458	0	6\n1	6\n2	3\n3	6...	

					YrSold	SaleType	SaleCondition
0	0	2010\n1	2010\n2	2010\n3	...	WD	Normal
1	0	2010\n1	2010\n2	2010\n3	...	WD	Normal
2	0	2010\n1	2010\n2	2010\n3	...	WD	Normal
3	0	2010\n1	2010\n2	2010\n3	...	WD	Normal
4	0	2010\n1	2010\n2	2010\n3	...	WD	Normal
...				
1454	0	2010\n1	2010\n2	2010\n3	...	WD	Normal
1455	0	2010\n1	2010\n2	2010\n3	...	WD	Abnorml
1456	0	2010\n1	2010\n2	2010\n3	...	WD	Abnorml
1457	0	2010\n1	2010\n2	2010\n3	...	WD	Normal
1458	0	2010\n1	2010\n2	2010\n3	...	WD	Normal

[1380 rows x 69 columns]

6.0.3 Checking skewness

```
[64]: #Choosing numerical columns
test_numerical_columns = test_df.select_dtypes(include = np.number).columns

for i, col in enumerate(test_numerical_columns):
    skewness = test_df[col].skew()
    kurt = test_df[col].kurtosis()
    print('Skewness of', col, 'is', skewness)
    print('Kurtosis of', col, 'is', kurt, '\n')
```

Skewness of LotArea is 3.115216613500925

Kurtosis of LotArea is 20.746548709480273

Skewness of YearBuilt is -0.5876566078696325

Kurtosis of YearBuilt is -0.579320618279759

Skewness of YearRemodAdd is -0.39990598881237394

Kurtosis of YearRemodAdd is -1.4125856813078284

Skewness of MasVnrArea is 2.549568580084032

Kurtosis of MasVnrArea is 8.479206132512642

Skewness of BsmtFinSF1 is 1.1663296764605604

Kurtosis of BsmtFinSF1 is 2.673291224416369

Skewness of BsmtFinSF2 is 4.0429539042374705

Kurtosis of BsmtFinSF2 is 17.68216539285966

Skewness of BsmtUnfSF is 0.9199232287642486

Kurtosis of BsmtUnfSF is 0.33244656805494754

Skewness of TotalBsmtSF is 0.8050662652894462

Kurtosis of TotalBsmtSF is 5.173387914148094

Skewness of 1stFlrSF is 1.558194572983426

Kurtosis of 1stFlrSF is 8.053863335994567

Skewness of 2ndFlrSF is 0.9128826344911904

Kurtosis of 2ndFlrSF is -0.27544098409544926

Skewness of LowQualFinSF is 16.167254030990343

Kurtosis of LowQualFinSF is 308.67690650153196

Skewness of GrLivArea is 1.1304024140503506

Kurtosis of GrLivArea is 2.9203451470914166

Skewness of BsmtFullBath is 0.6518652012126019
Kurtosis of BsmtFullBath is -0.6425404276091546

Skewness of BsmtHalfBath is 3.782975779434155
Kurtosis of BsmtHalfBath is 13.575839867420619

Skewness of FullBath is 0.2958386394763122
Kurtosis of FullBath is -0.232339259823799

Skewness of HalfBath is 0.7147275486062835
Kurtosis of HalfBath is -0.9887003982031177

Skewness of BedroomAbvGr is 0.43662327937443507
Kurtosis of BedroomAbvGr is 1.6859653302114266

Skewness of KitchenAbvGr is 4.079055022333038
Kurtosis of KitchenAbvGr is 17.471691151777893

Skewness of TotRmsAbvGrd is 0.8425974464281701
Kurtosis of TotRmsAbvGrd is 1.5225956466723316

Skewness of Fireplaces is 0.8198582704555165
Kurtosis of Fireplaces is 0.3873261355352815

Skewness of GarageCars is -0.10999345834943444
Kurtosis of GarageCars is 0.24648590262300463

Skewness of GarageArea is 0.29629030596790457
Kurtosis of GarageArea is 0.9610046428994408

Skewness of WoodDeckSF is 2.130759950574929
Kurtosis of WoodDeckSF is 10.249278055421557

Skewness of OpenPorchSF is 2.6877788503224047
Kurtosis of OpenPorchSF is 13.010835704951141

Skewness of EnclosedPorch is 4.669172309715744
Kurtosis of EnclosedPorch is 40.129017354477156

Skewness of 3SsnPorch is 12.524215926005567
Kurtosis of 3SsnPorch is 170.20011269208473

Skewness of ScreenPorch is 3.7882443811743776
Kurtosis of ScreenPorch is 17.239542060670402

Skewness of PoolArea is 20.196887591116827
Kurtosis of PoolArea is 445.66110156800477

Skewness of MiscVal is 20.075188353860344
Kurtosis of MiscVal is 471.51738779462795

```
[65]: # Selecting features with high skew
test_high_skew = high_skew(test_df, test_numerical_columns)
```

```
[66]: # Performing boxcox transformation for features with high skew
high_skewed_features_test_df_box = scipy.special.
    ↳ boxcox1p(test_df[test_high_skew], 0.5)
```

```
[67]: for i, col in enumerate(test_high_skew):
        skewness = high_skewed_features_test_df_box[col].skew()
        print('Skewness of', col, 'is', skewness)
```

Skewness of LotArea is 0.7363390828698114
Skewness of MasVnrArea is 1.149356822074124
Skewness of BsmtFinSF1 is 0.04578881428414585
Skewness of BsmtFinSF2 is 2.980868196365515
Skewness of 1stFlrSF is 0.6486423873019207
Skewness of LowQualFinSF is 12.029002531472997
Skewness of GrLivArea is 0.538179736134212
Skewness of BsmtHalfBath is 3.6894768259301016
Skewness of KitchenAbvGr is 3.757919518725453
Skewness of WoodDeckSF is 0.5858574655441303
Skewness of OpenPorchSF is 0.747153597002874
Skewness of EnclosedPorch is 2.3182627997359506
Skewness of 3SsnPorch is 11.085387021724689
Skewness of ScreenPorch is 2.9726015541097697
Skewness of PoolArea is 17.173516060992487
Skewness of MiscVal is 9.611681553124736

Clearly skewness has been reduced.

```
[68]: # Replacing features with high skew with transformed values
test_df[test_high_skew] = high_skewed_features_test_df_box
```

6.0.4 Categorical Variables

Lets encode all of the categorical variables.

```
[69]: test_df = pd.get_dummies(test_df)
test_df
```

```
[69]:
```

	LotArea	YearBuilt	YearRemodAdd	MasVnrArea	BsmtFinSF1	BsmtFinSF2	\
0	213.620036	1961	1961	0.000000	41.312816	22.083189	
1	236.897468	1958	1958	18.880613	58.794737	0.000000	
2	233.210544	1997	1998	0.000000	54.284989	0.000000	

3	197.789890	1998	1998	7.165151	47.112117	0.000000
4	139.506184	1992	1992	0.000000	30.496154	0.000000
...
1454	86.022724	1970	1970	0.000000	0.000000	0.000000
1455	85.063195	1970	1970	0.000000	29.811947	0.000000
1456	280.849783	1960	1996	0.000000	68.000000	0.000000
1457	202.372209	1992	1992	0.000000	34.769553	0.000000
1458	194.244745	1993	1994	17.493589	53.099909	0.000000

	BsmtUnfSF	TotalBsmtSF	1stFlrSF	2ndFlrSF	...	SaleType_ConLw	\
0	270.0	882.0	57.899917	0	...	0	
1	406.0	1329.0	70.938330	0	...	0	
2	137.0	928.0	58.959003	701	...	0	
3	324.0	926.0	58.893349	678	...	0	
4	1017.0	1280.0	69.582121	0	...	0	
...	
1454	546.0	546.0	44.776062	546	...	0	
1455	294.0	546.0	44.776062	546	...	0	
1456	0.0	1224.0	68.000000	0	...	0	
1457	575.0	912.0	60.321746	0	...	0	
1458	238.0	996.0	61.150614	1004	...	0	

	SaleType_New	SaleType_Oth	SaleType_WD	SaleCondition_Abnorml	\
0	0	0	1	0	
1	0	0	1	0	
2	0	0	1	0	
3	0	0	1	0	
4	0	0	1	0	
...	
1454	0	0	1	0	
1455	0	0	1	1	
1456	0	0	1	1	
1457	0	0	1	0	
1458	0	0	1	0	

	SaleCondition_AdjLand	SaleCondition_Alloca	SaleCondition_Family	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	
...	
1454	0	0	0	
1455	0	0	0	
1456	0	0	0	
1457	0	0	0	
1458	0	0	0	

	SaleCondition_Normal	SaleCondition_Partial
0	1	0
1	1	0
2	1	0
3	1	0
4	1	0
...
1454	1	0
1455	0	0
1456	0	0
1457	1	0
1458	1	0

[1459 rows x 254 columns]

6.0.5 Comparing Test and Train datasets

Let me double check the test and train sets are similar so the test set can pass through the model.

```
[70]: print(train_df.shape)
      print(test_df.shape)
```

(1362, 255)

(1459, 254)

I am going to examine the column names to see if there are any columns which are in the test set but not in the training set and I will remove these columns

```
[71]: list(test_df.columns)
```

```
[71]: ['LotArea',
      'YearBuilt',
      'YearRemodAdd',
      'MasVnrArea',
      'BsmtFinSF1',
      'BsmtFinSF2',
      'BsmtUnfSF',
      'TotalBsmtSF',
      '1stFlrSF',
      '2ndFlrSF',
      'LowQualFinSF',
      'GrLivArea',
      'BsmtFullBath',
      'BsmtHalfBath',
      'FullBath',
      'HalfBath',
      'BedroomAbvGr',
```

```

'KitchenAbvGr',
'TotRmsAbvGrd',
'Fireplaces',
'GarageCars',
'GarageArea',
'WoodDeckSF',
'OpenPorchSF',
'EnclosedPorch',
'3SsnPorch',
'ScreenPorch',
'PoolArea',
'MiscVal',
'MSSubClass_0      20\n1      20\n2      60\n3      60\n4      120\n
... \n1454      160\n1455      160\n1456      20\n1457      85\n1458      60\nName:
MSSubClass, Length: 1459, dtype: int64',
'MSZoning_0',
'MSZoning_C (all)',
'MSZoning_FV',
'MSZoning_RH',
'MSZoning_RL',
'MSZoning_RM',
'Street_Grvl',
'Street_Pave',
'LotShape_IR1',
'LotShape_IR2',
'LotShape_IR3',
'LotShape_Reg',
'LandContour_Bnk',
'LandContour_HLS',
'LandContour_Low',
'LandContour_Lvl',
'Utilities_0',
'Utilities_AllPub',
'LotConfig_Corner',
'LotConfig_CulDSac',
'LotConfig_FR2',
'LotConfig_FR3',
'LotConfig_Inside',
'LandSlope_Gtl',
'LandSlope_Mod',
'LandSlope_Sev',
'Neighborhood_Blmngtn',
'Neighborhood_Blueste',
'Neighborhood_BrDale',
'Neighborhood_BrkSide',
'Neighborhood_ClearCr',
'Neighborhood_CollgCr',

```

'Neighborhood_Crawfor',
 'Neighborhood_Edwards',
 'Neighborhood_Gilbert',
 'Neighborhood_IDOTRR',
 'Neighborhood_MeadowV',
 'Neighborhood_Mitchel',
 'Neighborhood_NAmes',
 'Neighborhood_NPkVill',
 'Neighborhood_NWAmes',
 'Neighborhood_NoRidge',
 'Neighborhood_NridgHt',
 'Neighborhood_OldTown',
 'Neighborhood_SWISU',
 'Neighborhood_Sawyer',
 'Neighborhood_SawyerW',
 'Neighborhood_Somerst',
 'Neighborhood_StoneBr',
 'Neighborhood_Timber',
 'Neighborhood_Veenker',
 'Condition1_Artery',
 'Condition1_Feedr',
 'Condition1_Norm',
 'Condition1_PosA',
 'Condition1_PosN',
 'Condition1_RRAe',
 'Condition1_RRAAn',
 'Condition1_RRNe',
 'Condition1_RRNn',
 'Condition2_Artery',
 'Condition2_Feedr',
 'Condition2_Norm',
 'Condition2_PosA',
 'Condition2_PosN',
 'BldgType_1Fam',
 'BldgType_2fmCon',
 'BldgType_Duplex',
 'BldgType_Twnhs',
 'BldgType_TwnhsE',
 'HouseStyle_1.5Fin',
 'HouseStyle_1.5Unf',
 'HouseStyle_1Story',
 'HouseStyle_2.5Unf',
 'HouseStyle_2Story',
 'HouseStyle_SFoyer',
 'HouseStyle_SLvl',
 'OverallQual_0 5\n1 6\n2 5\n3 6\n4 8\n
 ..\n1454 4\n1455 4\n1456 5\n1457 5\n1458 7\nName: OverallQual,


```

Length: 1459, dtype: int64',
'OverallCond_0      6\n1      6\n2      5\n3      6\n4      5\n
..\n1454      7\n1455      5\n1456      7\n1457      5\n1458      5\nName: OverallCond,
Length: 1459, dtype: int64',
'RoofStyle_Flat',
'RoofStyle_Gable',
'RoofStyle_Gambrel',
'RoofStyle_Hip',
'RoofStyle_Mansard',
'RoofStyle_Shed',
'RoofMatl_CompShg',
'RoofMatl_Tar&Grv',
'RoofMatl_WdShake',
'RoofMatl_WdShngl',
'Exterior1st_0',
'Exterior1st_AsbShng',
'Exterior1st_AsphShn',
'Exterior1st_BrkComm',
'Exterior1st_BrkFace',
'Exterior1st_CBlock',
'Exterior1st_CemntBd',
'Exterior1st_HdBoard',
'Exterior1st_MetalSd',
'Exterior1st_Plywood',
'Exterior1st_Stucco',
'Exterior1st_VinylSd',
'Exterior1st_Wd Sdng',
'Exterior1st_WdShing',
'Exterior2nd_0',
'Exterior2nd_AsbShng',
'Exterior2nd_AsphShn',
'Exterior2nd_Brk Cmn',
'Exterior2nd_BrkFace',
'Exterior2nd_CBlock',
'Exterior2nd_CmentBd',
'Exterior2nd_HdBoard',
'Exterior2nd_ImStucc',
'Exterior2nd_MetalSd',
'Exterior2nd_Plywood',
'Exterior2nd_Stone',
'Exterior2nd_Stucco',
'Exterior2nd_VinylSd',
'Exterior2nd_Wd Sdng',
'Exterior2nd_Wd Shng',
'MasVnrType_0',
'MasVnrType_BrkCmn',
'MasVnrType_BrkFace',

```

'MasVnrType_None',
'MasVnrType_Stone',
'ExterQual_Ex',
'ExterQual_Fa',
'ExterQual_Gd',
'ExterQual_TA',
'ExterCond_Ex',
'ExterCond_Fa',
'ExterCond_Gd',
'ExterCond_Po',
'ExterCond_TA',
'Foundation_BrkTil',
'Foundation_CBlock',
'Foundation_PConc',
'Foundation_Slab',
'Foundation_Stone',
'Foundation_Wood',
'BsmtCond_0',
'BsmtCond_Fa',
'BsmtCond_Gd',
'BsmtCond_Po',
'BsmtCond_TA',
'BsmtExposure_0',
'BsmtExposure_Av',
'BsmtExposure_Gd',
'BsmtExposure_Mn',
'BsmtExposure_No',
'BsmtFinType1_0',
'BsmtFinType1_ALQ',
'BsmtFinType1_BLQ',
'BsmtFinType1_GLQ',
'BsmtFinType1_LwQ',
'BsmtFinType1_Rec',
'BsmtFinType1_Unf',
'BsmtFinType2_0',
'BsmtFinType2_ALQ',
'BsmtFinType2_BLQ',
'BsmtFinType2_GLQ',
'BsmtFinType2_LwQ',
'BsmtFinType2_Rec',
'BsmtFinType2_Unf',
'Heating_GasA',
'Heating_GasW',
'Heating_Grav',
'Heating_Wall',
'HeatingQC_Ex',
'HeatingQC_Fa',

```

'HeatingQC_Gd',
'HeatingQC_Po',
'HeatingQC_TA',
'CentralAir_N',
'CentralAir_Y',
'Electrical_FuseA',
'Electrical_FuseF',
'Electrical_FuseP',
'Electrical_SBrkr',
'KitchenQual_0',
'KitchenQual_Ex',
'KitchenQual_Fa',
'KitchenQual_Gd',
'KitchenQual_TA',
'Functional_0',
'Functional_Maj1',
'Functional_Maj2',
'Functional_Min1',
'Functional_Min2',
'Functional_Mod',
'Functional_Sev',
'Functional_Typ',
'GarageType_0',
'GarageType_2Types',
'GarageType_Attchd',
'GarageType_Basment',
'GarageType_BuiltIn',
'GarageType_CarPort',
'GarageType_Detchd',
'GarageQual_0',
'GarageQual_Fa',
'GarageQual_Gd',
'GarageQual_Po',
'GarageQual_TA',
'PavedDrive_N',
'PavedDrive_P',
'PavedDrive_Y',
'MoSold_0      6\n1      6\n2      3\n3      6\n4      1\n
..\n1454      6\n1455      4\n1456      9\n1457      7\n1458      11\nName: MoSold,
Length: 1459, dtype: int64',
'YrSold_0      2010\n1      2010\n2      2010\n3      2010\n4      2010\n
... \n1454      2006\n1455      2006\n1456      2006\n1457      2006\n1458
2006\nName: YrSold, Length: 1459, dtype: int64',
'SaleType_0',
'SaleType_COD',
'SaleType_CWD',
'SaleType_Con',

```

```

'SaleType_ConLD',
'SaleType_ConLI',
'SaleType_ConLw',
'SaleType_New',
'SaleType_Oth',
'SaleType_WD',
'SaleCondition_Abnorml',
'SaleCondition_AdjLand',
'SaleCondition_Alloca',
'SaleCondition_Family',
'SaleCondition_Normal',
'SaleCondition_Partial']

```

```
[72]: list(train_df.columns)
```

```

[72]: ['Id',
      'LotArea',
      'YearBuilt',
      'YearRemodAdd',
      'MasVnrArea',
      'BsmtFinSF1',
      'BsmtFinSF2',
      'BsmtUnfSF',
      'TotalBsmtSF',
      '1stFlrSF',
      '2ndFlrSF',
      'LowQualFinSF',
      'GrLivArea',
      'BsmtFullBath',
      'BsmtHalfBath',
      'FullBath',
      'HalfBath',
      'BedroomAbvGr',
      'KitchenAbvGr',
      'TotRmsAbvGrd',
      'Fireplaces',
      'GarageCars',
      'GarageArea',
      'WoodDeckSF',
      'OpenPorchSF',
      'EnclosedPorch',
      '3SsnPorch',
      'ScreenPorch',
      'PoolArea',
      'MiscVal',
      'SalePrice',
      'MSSubClass_0      60\n1      20\n2      60\n3      70\n4      60\n

```

```

..\n1455    60\n1456    20\n1457    70\n1458    20\n1459    20\nName:
MSSubClass, Length: 1362, dtype: int64',
'MSZoning_C (all)',
'MSZoning_FV',
'MSZoning_RH',
'MSZoning_RL',
'MSZoning_RM',
'Street_Grvl',
'Street_Pave',
'LotShape_IR1',
'LotShape_IR2',
'LotShape_IR3',
'LotShape_Reg',
'LandContour_Bnk',
'LandContour_HLS',
'LandContour_Low',
'LandContour_Lvl',
'Utilities_AllPub',
'Utilities_NoSeWa',
'LotConfig_Corner',
'LotConfig_CulDSac',
'LotConfig_FR2',
'LotConfig_FR3',
'LotConfig_Inside',
'LandSlope_Gtl',
'LandSlope_Mod',
'LandSlope_Sev',
'Neighborhood_Blmngtn',
'Neighborhood_Blueste',
'Neighborhood_BrDale',
'Neighborhood_BrkSide',
'Neighborhood_ClearCr',
'Neighborhood_CollgCr',
'Neighborhood_Crawfor',
'Neighborhood_Edwards',
'Neighborhood_Gilbert',
'Neighborhood_IDOTRR',
'Neighborhood_MeadowV',
'Neighborhood_Mitchel',
'Neighborhood_NAmes',
'Neighborhood_NPkVill',
'Neighborhood_NWAmes',
'Neighborhood_NoRidge',
'Neighborhood_NridgHt',
'Neighborhood_OldTown',
'Neighborhood_SWISU',
'Neighborhood_Sawyer',

```

```

'Neighborhood_SawyerW',
'Neighborhood_Somerst',
'Neighborhood_StoneBr',
'Neighborhood_Timber',
'Neighborhood_Veenker',
'Condition1_Artery',
'Condition1_Feedr',
'Condition1_Norm',
'Condition1_PosA',
'Condition1_PosN',
'Condition1_RRAe',
'Condition1_RRAn',
'Condition1_RRNe',
'Condition1_RRNn',
'Condition2_Artery',
'Condition2_Feedr',
'Condition2_Norm',
'Condition2_PosN',
'Condition2_RRAe',
'Condition2_RRAn',
'Condition2_RRNn',
'BldgType_1Fam',
'BldgType_2fmCon',
'BldgType_Duplex',
'BldgType_Twnhs',
'BldgType_TwnhsE',
'HouseStyle_1.5Fin',
'HouseStyle_1.5Unf',
'HouseStyle_1Story',
'HouseStyle_2.5Fin',
'HouseStyle_2.5Unf',
'HouseStyle_2Story',
'HouseStyle_SFoyer',
'HouseStyle_SLvl',
'OverallQual_0      7\n1      6\n2      7\n3      7\n4      8\n
..\n1455      6\n1456      6\n1457      7\n1458      5\n1459      5\nName: OverallQual,
Length: 1362, dtype: int64',
'OverallCond_0      5\n1      8\n2      5\n3      5\n4      5\n
..\n1455      5\n1456      6\n1457      9\n1458      6\n1459      6\nName: OverallCond,
Length: 1362, dtype: int64',
'RoofStyle_Flat',
'RoofStyle_Gable',
'RoofStyle_Gambrel',
'RoofStyle_Hip',
'RoofStyle_Mansard',
'RoofStyle_Shed',
'RoofMatl_CompShg',

```

'RoofMatl_Membran',
 'RoofMatl_Metal',
 'RoofMatl_Roll',
 'RoofMatl_Tar&Grv',
 'RoofMatl_WdShake',
 'RoofMatl_WdShngl',
 'Exterior1st_AsbShng',
 'Exterior1st_BrkFace',
 'Exterior1st_CBlock',
 'Exterior1st_CemntBd',
 'Exterior1st_HdBoard',
 'Exterior1st_ImStucc',
 'Exterior1st_MetalSd',
 'Exterior1st_Plywood',
 'Exterior1st_Stone',
 'Exterior1st_Stucco',
 'Exterior1st_VinylSd',
 'Exterior1st_Wd Sdng',
 'Exterior1st_WdShing',
 'Exterior2nd_AsbShng',
 'Exterior2nd_AsphShn',
 'Exterior2nd_Brk Cmn',
 'Exterior2nd_BrkFace',
 'Exterior2nd_CBlock',
 'Exterior2nd_CmentBd',
 'Exterior2nd_HdBoard',
 'Exterior2nd_ImStucc',
 'Exterior2nd_MetalSd',
 'Exterior2nd_Other',
 'Exterior2nd_Plywood',
 'Exterior2nd_Stone',
 'Exterior2nd_Stucco',
 'Exterior2nd_VinylSd',
 'Exterior2nd_Wd Sdng',
 'Exterior2nd_Wd Shng',
 'MasVnrType_BrkCmn',
 'MasVnrType_BrkFace',
 'MasVnrType_None',
 'MasVnrType_Stone',
 'ExterQual_Ex',
 'ExterQual_Fa',
 'ExterQual_Gd',
 'ExterQual_TA',
 'ExterCond_Ex',
 'ExterCond_Fa',
 'ExterCond_Gd',
 'ExterCond_Po',

'ExterCond_TA',
'Foundation_BrkTil',
'Foundation_CBlock',
'Foundation_PConc',
'Foundation_Stone',
'Foundation_Wood',
'BsmtQual_Ex',
'BsmtQual_Fa',
'BsmtQual_Gd',
'BsmtQual_TA',
'BsmtCond_Fa',
'BsmtCond_Gd',
'BsmtCond_Po',
'BsmtCond_TA',
'BsmtExposure_Av',
'BsmtExposure_Gd',
'BsmtExposure_Mn',
'BsmtExposure_No',
'BsmtFinType1_ALQ',
'BsmtFinType1_BLQ',
'BsmtFinType1_GLQ',
'BsmtFinType1_LwQ',
'BsmtFinType1_Rec',
'BsmtFinType1_Unf',
'BsmtFinType2_ALQ',
'BsmtFinType2_BLQ',
'BsmtFinType2_GLQ',
'BsmtFinType2_LwQ',
'BsmtFinType2_Rec',
'BsmtFinType2_Unf',
'Heating_GasA',
'Heating_GasW',
'Heating_Grav',
'Heating_OthW',
'HeatingQC_Ex',
'HeatingQC_Fa',
'HeatingQC_Gd',
'HeatingQC_Po',
'HeatingQC_TA',
'CentralAir_N',
'CentralAir_Y',
'Electrical_FuseA',
'Electrical_FuseF',
'Electrical_FuseP',
'Electrical_Mix',
'Electrical_SBrkr',
'KitchenQual_Ex',


```

'KitchenQual_Fa',
'KitchenQual_Gd',
'KitchenQual_TA',
'Functional_Maj1',
'Functional_Maj2',
'Functional_Min1',
'Functional_Min2',
'Functional_Mod',
'Functional_Sev',
'Functional_Typ',
'GarageType_2Types',
'GarageType_Attchd',
'GarageType_Basment',
'GarageType_BuiltIn',
'GarageType_CarPort',
'GarageType_Detchd',
'GarageQual_Ex',
'GarageQual_Fa',
'GarageQual_Gd',
'GarageQual_Po',
'GarageQual_TA',
'PavedDrive_N',
'PavedDrive_P',
'PavedDrive_Y',
'MoSold_0      2\n1      5\n2      9\n3      2\n4      12\n
..\n1455      8\n1456      2\n1457      5\n1458      4\n1459      6\nName: MoSold,
Length: 1362, dtype: int64',
'YrSold_0      2008\n1      2007\n2      2008\n3      2006\n4      2008\n
... \n1455      2007\n1456      2010\n1457      2010\n1458      2010\n1459
2008\nName: YrSold, Length: 1362, dtype: int64',
'SaleType_COD',
'SaleType_CWD',
'SaleType_Con',
'SaleType_ConLD',
'SaleType_ConLI',
'SaleType_ConLw',
'SaleType_New',
'SaleType_Oth',
'SaleType_WD',
'SaleCondition_Abnorml',
'SaleCondition_AdjLand',
'SaleCondition_Alloca',
'SaleCondition_Family',
'SaleCondition_Normal',
'SaleCondition_Partial']

```

'Utilities_0' is a feature which should be removed

```
[73]: test_df = test_df.drop(columns = 'Utilities_0')
```

Now the test set is ready to pass through my models

7 Modelling

First lets split our data up into the features we want to learn from and the features we want to test on.

```
[74]: #Dropping the column Id as it is just a label for the row. We can use inbuilt  
      →Python index  
train_df = train_df.drop(columns = 'Id')
```

```
[75]: #These are the features which our model will use to make predictions about  
      →SalePrice  
X = train_df.drop(columns = 'SalePrice')
```

```
[76]: #Our target variable  
y = train_df['SalePrice']
```

I will now split X and y into train and test splits using StratifiedKFold which preserves the distribution of the data and uses all the data to train and test models.

```
[77]: skf = StratifiedKFold(n_splits=10, shuffle=True, random_state=42)  
train_df_split = skf.get_n_splits(X,y)
```

7.0.1 Linear Regression

I will first fit a simple linear regression model and find the model accuracy using the error metric `cross_val_score`. For more information on how the `LinearRegression` model operates: https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LinearRegression.html

```
[78]: lr = LinearRegression()  
  
      #Calculating cross_val_score for the 10 folds  
train_lr_scores = (-1) * cross_val_score(lr, X, y, cv = train_df_split, scoring_  
      →= 'neg_mean_absolute_error')  
  
print(train_lr_scores)
```

```
[0.07766194 0.08720395 0.08578021 0.08860303 0.09179059 0.07992134  
0.07895251 0.07925374 0.08040423 0.08601455]
```

```
[79]: #Calculating the mean of these scores  
train_lr_scores_mean = np.mean(train_lr_scores)  
print(train_lr_scores_mean)
```

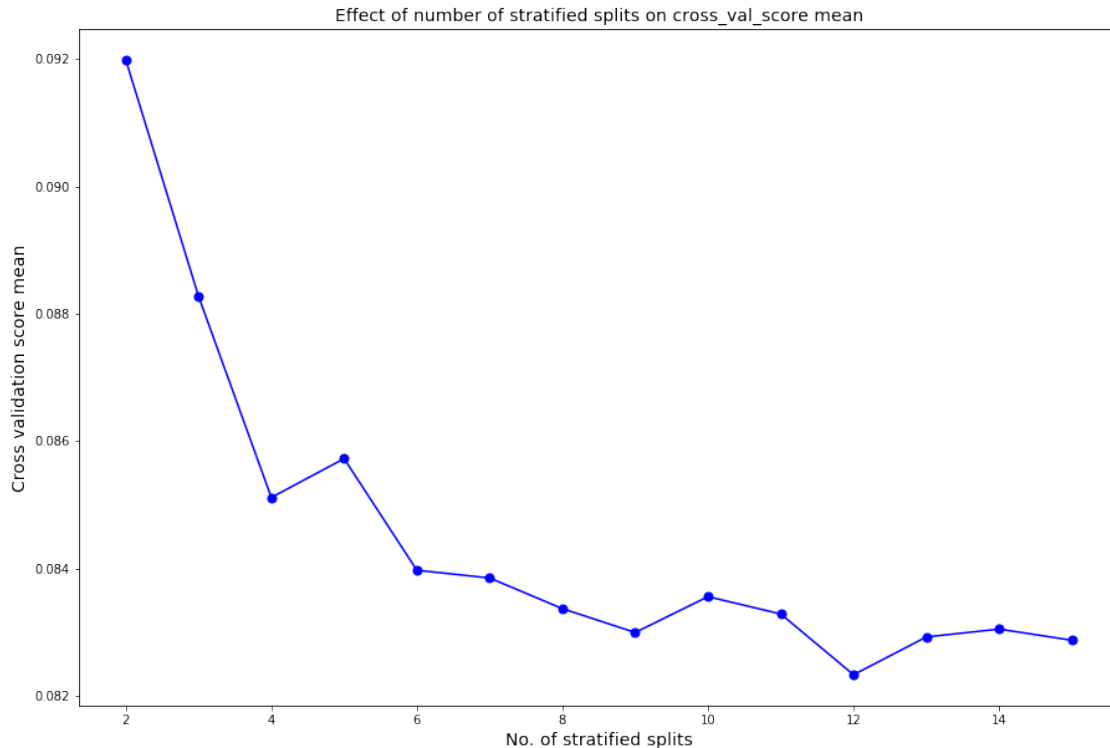
0.08355861009621796

An MAE of 0.084! This is quite good. However we just picked a random number of folds. Lets see how this value varies with the number of folds.

```
[80]: def best_cross_val_score(model, data_X, data_y, splits):  
  
    """  
    This function will give the best cross validation score from a range of  
    →stratified splits  
  
    Input: machine learning model, input features, target features, no. of  
    →desired splits  
  
    Output: Best mean cross validation score, plot of mean cross validation  
    →score with no. of splits  
  
    """  
  
    data_mean_array = []  
    split_array = []  
  
    for i, split in enumerate(range(2, splits+1)):  
  
        skf_loop = StratifiedKFold(n_splits=split, shuffle=True, random_state=42)  
        data_split = skf_loop.get_n_splits(data_X, data_y)  
        data_scores = (-1) * cross_val_score(model, data_X, data_y, cv =  
    →data_split, scoring = 'neg_mean_absolute_error')  
        data_scores_mean = np.mean(data_scores)  
        data_mean_array.append(data_scores_mean)  
        split_array.append(split)  
  
    print('The most accurate mean cross validation score is',  
    →max(data_mean_array))  
  
    plt.figure(figsize = (15,10))  
    plt.plot(split_array, data_mean_array, 'b', marker='.', markersize=14)  
    plt.title("Effect of number of stratified splits on cross_val_score mean",  
    →fontsize = 14)  
    plt.xlabel("No. of stratified splits", fontsize = 14)  
    plt.ylabel("Cross validation score mean", fontsize = 14)  
    plt.show()
```

```
[81]: #Finding the number of splits which results in the best cross validation score  
best_cross_val_score(lr, X, y, 15)
```

The most accurate mean cross validation score is 0.09198970379208748



So we can see that as the number of folds the CVS mean increases. Therefore our accuracy increases. I used splits = 15 as any greater, it becomes computationally expensive to run the code whilst the accuracy does not improve by much. Therefore the 12 is the best number of stratified splits to use.

```
[82]: # Storing the most accurate model
skf = StratifiedKFold(n_splits=12, shuffle=True, random_state=42)
train_df_split = skf.get_n_splits(X,y)
```

The number of splits will be useful for when I start to hypertune parameters to build more accurate models. Now I will split my training data so that I can fit in a simple linear regression model.

```
[83]: # Splitting data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size = 0.2,
↳random_state = 42)
```

```
[84]: # Fitting model to training split data
lr.fit(X_train,y_train)

# Predicting SalePrice using model on test split data
y_pred_lr = lr.predict(X_test)

# Calculating mean absolute error between predicted values and actual values
```

```
print(mean_absolute_error(y_pred_lr, y_test))
```

0.08945059791046515

I achieve a low mean absolute error of 0.0895. This is impressive. However I think we can reduce this by finding a better model.

Now, let me use this model to predict on my test set

```
[85]: # Storing prediction on test set
lr_pred = lr.predict(test_df)
```

7.0.2 Ridge Regression

I will now use a ridge regression model which should hopefully manage some of the over-fitting of a standard ordinary least squares model. For more information: https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.Ridge.html

```
[86]: rr = Ridge(alpha = 1.0, random_state=42)
```

```
[87]: rr.fit(X_train, y_train)

y_pred_rr = rr.predict(X_test)

print(mean_absolute_error(y_pred_rr, y_test))
```

0.0852248591197156

Now we have used the ridge regression model we have reduced the mean absolute error. Although the reduction is minimal.

However, we have only done this for 1 alpha value, let me do this for a range of alpha values. I will use the GridSearchCV function for this. GridSearchCV allows me to do an exhaustive search over a range of parameter values. For more information: https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV.html

```
[88]: # Some warnings were appearing so I will ignore them for now
from sklearn.model_selection import GridSearchCV
import warnings

warnings.filterwarnings('ignore')
```

As I am tuning hyperparameters I will use a cross validation splitting strategy using my best stratified k folds splitting strategy found previously.

```
[89]: # Defining range of parameters to use
parameters = {'alpha': [1e-15, 1e-10, 1e-8, 1e-4, 1e-3, 1e-2, 1, 5, 10, 20]}

ridge_regressor = GridSearchCV(rr, parameters, scoring =_
    ↪ 'neg_mean_absolute_error', cv = train_df_split)
```

```
ridge_regressor.fit(X, y)
```

```
[89]: GridSearchCV(cv=12, error_score='raise-deprecating',
                  estimator=Ridge(alpha=1.0, copy_X=True, fit_intercept=True,
                                  max_iter=None, normalize=False, random_state=42,
                                  solver='auto', tol=0.001),
                  iid='warn', n_jobs=None,
                  param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.0001, 0.001, 0.01, 1,
                                         5, 10, 20]},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                  scoring='neg_mean_absolute_error', verbose=0)
```

```
[90]: print('The best parameter for', ridge_regressor.best_params_)
      print('The best mean absolute error:', (-1) * ridge_regressor.best_score_)
```

The best parameter for {'alpha': 5}

The best mean absolute error: 0.07871660662991475

Clearly, now the mean absolute error has reduced by a much greater amount and we have found a better model. Therefore, I will use this model to make my predictions with the test data

```
[91]: rr_pred = ridge_regressor.predict(test_df)
```

7.0.3 Lasso Regression

Now I will use a lasso regression model to check whether anymore improvements can be made. For more information: https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.Lasso.html

```
[92]: lasso = Lasso(random_state = 42)

lasso_regressor = GridSearchCV(lasso, parameters, scoring =_
    → 'neg_mean_absolute_error', cv = train_df_split)

lasso_regressor.fit(X, y)
```

```
[92]: GridSearchCV(cv=12, error_score='raise-deprecating',
                  estimator=Lasso(alpha=1.0, copy_X=True, fit_intercept=True,
                                  max_iter=1000, normalize=False, positive=False,
                                  precompute=False, random_state=42,
                                  selection='cyclic', tol=0.0001, warm_start=False),
                  iid='warn', n_jobs=None,
                  param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.0001, 0.001, 0.01, 1,
                                         5, 10, 20]},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                  scoring='neg_mean_absolute_error', verbose=0)
```

```
[93]: print('The best parameter for', lasso_regressor.best_params_)
      print('The best mean absolute error:', (-1) * lasso_regressor.best_score_)
```

The best parameter for {'alpha': 0.0001}
 The best mean absolute error: 0.07947369397324025

The lasso regression model does not improve on the mean absolute error hence I will not make predictions using this model.

7.0.4 Elastic Net

Now I will use the Elastic Net model which uses both ridge regression and lasso regression to build a model which should help manage overfitting and underfitting. For documentation on this model: https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.ElasticNet.html

```
[94]: from sklearn.linear_model import ElasticNet
```

```
[95]: parameters_grid = {'alpha': [1e-15, 1e-10, 1e-8, 1e-4, 1e-3, 1e-2, 1, 5, 10, 20],
                        'l1_ratio': [0, 0.1, 1]}
```

```
[96]: elastic = ElasticNet(random_state = 42)

elastic_regressor = GridSearchCV(elastic, parameters_grid, scoring =_
    → 'neg_mean_absolute_error', cv = train_df_split)

elastic_regressor.fit(X, y)
```

```
[96]: GridSearchCV(cv=12, error_score='raise-deprecating',
                  estimator=ElasticNet(alpha=1.0, copy_X=True, fit_intercept=True,
                                       l1_ratio=0.5, max_iter=1000, normalize=False,
                                       positive=False, precompute=False,
                                       random_state=42, selection='cyclic',
                                       tol=0.0001, warm_start=False),
                  iid='warn', n_jobs=None,
                  param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.0001, 0.001, 0.01, 1,
                                       5, 10, 20],
                              'l1_ratio': [0, 0.1, 1]},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                  scoring='neg_mean_absolute_error', verbose=0)
```

```
[97]: print('The best parameter for', elastic_regressor.best_params_)
      print('The best mean absolute error:', (-1) * elastic_regressor.best_score_)
```

The best parameter for {'alpha': 0.01, 'l1_ratio': 0}
 The best mean absolute error: 0.07883602214880273

As we can see the mean absolute error does not improve. Also, this is a computationally expensive model to run for only a likely minimally improved model hence it will not be used.

7.0.5 Decision Tree Regressor

A decision tree aims to create a predictive model by learning decision rules which are inferred from the features of the dataset. For more information on the model: <https://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeRegressor.html> and <https://scikit-learn.org/stable/modules/tree.html#tree>

```
[98]: dt = DecisionTreeRegressor(criterion = 'mae', random_state = 42)

parameters_dt = {'max_depth': [3, 5],
                  'max_leaf_nodes': [3, 5]}

dt_regressor = GridSearchCV(dt, parameters_dt, cv = train_df_split, scoring = 'neg_mean_absolute_error')

dt_regressor.fit(X, y)
```

```
[98]: GridSearchCV(cv=12, error_score='raise-deprecating',
                  estimator=DecisionTreeRegressor(criterion='mae', max_depth=None,
                                                  max_features=None,
                                                  max_leaf_nodes=None,
                                                  min_impurity_decrease=0.0,
                                                  min_impurity_split=None,
                                                  min_samples_leaf=1,
                                                  min_samples_split=2,
                                                  min_weight_fraction_leaf=0.0,
                                                  presort=False, random_state=42,
                                                  splitter='best'),
                  iid='warn', n_jobs=None,
                  param_grid={'max_depth': [3, 5], 'max_leaf_nodes': [3, 5]},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                  scoring='neg_mean_absolute_error', verbose=0)
```

```
[99]: print('The best parameter for', dt_regressor.best_params_)
      print('The best mean absolute error:', (-1) * dt_regressor.best_score_)
```

```
The best parameter for {'max_depth': 3, 'max_leaf_nodes': 5}
The best mean absolute error: 0.17712525691570757
```

If we add more parameters to further hypertune this model, it is very computationally expensive and runs very slowly. Also, the mean absolute error is a lot greater than the ones achieved using ridge regression hence this is not an ideal model to use. Let's take a deeper look as to why this model has not worked as well.

In the plot below, I do not use the same splitting of the training data as I did in my calculation of the mean absolute error. This is because I am not concerned so much about the accuracy of the following model, but to demonstrate where this type of model fails when working with the training data I use.


```
[100]: from sklearn import tree

dt_plot = DecisionTreeRegressor(criterion = 'mae', random_state = 42, max_depth_
    ↳ 3, max_leaf_nodes = 5)

dt_plot = dt_plot.fit(X, y)

from sklearn.tree import export_graphviz

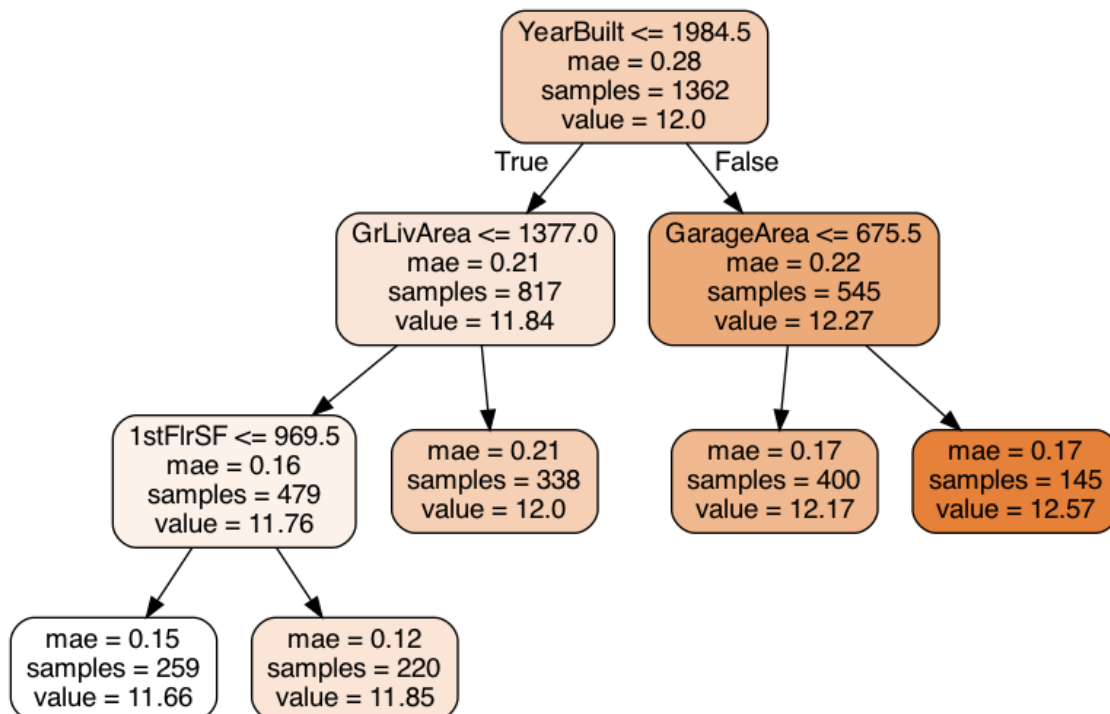
# Export as dot file
export_graphviz(dt_plot, out_file='tree.dot',
                feature_names = X.columns,
                rounded = True, proportion = False,
                precision = 2, filled = True)

import pydot

(graph,) = pydot.graph_from_dot_file(r'tree.dot')
graph.write_png(r'tree.png')

# Display in jupyter notebook
from IPython.display import Image
Image(filename = r'tree.png')
```

[100]:



As we can see from the decision tree plot above, this is not a great model for our dataset as we have lots of features which makes the decision tree computationally expensive to run. Additionally, if we do use a decision tree, I will only use a small number of features, making my model prone to underfitting. Therefore, this model leads to a high mean absolute error.

7.0.6 Random Forest Regressor

Now I will use a random forest tree, as this model uses many trees and makes predictions by averaging the predictions of each component tree.

This will help use as many features as possible when training a model hence managing underfitting and overfitting. For more information:

<https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestRegressor.html>

Instead of using GridSearchCV, I use RandomizedSearchCV as I want to do a random search on the hyperparameters used to reduced computational cost. For more

information: https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.RandomizedSearchCV.html

```
[101]: from sklearn.model_selection import RandomizedSearchCV

rf = RandomForestRegressor(criterion = 'mae', random_state = 42)

parameters_rf = {'n_estimators': [10],
                  'max_depth': [3]}

rf_regressor = RandomizedSearchCV(rf, parameters_rf, cv = train_df_split,
    random_state = 42, n_jobs = -1,
    scoring = 'neg_mean_absolute_error')
```

```
[102]: rf_regressor.fit(X, y)
```

```
[102]: RandomizedSearchCV(cv=12, error_score='raise-deprecating',
    estimator=RandomForestRegressor(bootstrap=True,
    criterion='mae',
    max_depth=None,
    max_features='auto',
    max_leaf_nodes=None,
    min_impurity_decrease=0.0,
    min_impurity_split=None,
    min_samples_leaf=1,
    min_samples_split=2,
    min_weight_fraction_leaf=0.0,
    n_estimators='warn',
    n_jobs=None, oob_score=False,
    random_state=42, verbose=0,
    warm_start=False),
    iid='warn', n_iter=10, n_jobs=-1,
    param_distributions={'max_depth': [3], 'n_estimators': [10]},
```

```
pre_dispatch='2*n_jobs', random_state=42, refit=True,
return_train_score=False, scoring='neg_mean_absolute_error',
verbose=0)
```

```
[103]: print('The best parameter for', rf_regressor.best_params_)
print('The best mean absolute error:', (-1) * rf_regressor.best_score_)
```

The best parameter for {'n_estimators': 10, 'max_depth': 3}
The best mean absolute error: 0.13293277300624817

Again, we have a high mean absolute error with a high computational cost model
hence this is not an ideal model to use

7.0.7 XGBoost Regressor

I will now use an aggressive gradient boosting regressor to iterate
over many random models to find the best model. For more information:
https://xgboost.readthedocs.io/en/latest/python/python_api.html#module-xgboost.sklearn

```
[104]: xgb = XGBRegressor(random_state = 42)

parameters_xgb = {'n_estimators': [10, 50, 100],
                  'learning_rate': [0.01, 0.025, 0.05]}

xgb_regressor = RandomizedSearchCV(xgb, parameters_xgb, cv = train_df_split,
→random_state = 42,
                                scoring = 'neg_mean_absolute_error')
```

```
[105]: xgb_regressor.fit(X, y)
```

```
[105]: RandomizedSearchCV(cv=12, error_score='raise-deprecating',
                        estimator=XGBRegressor(base_score=None, booster=None,
                                                colsample_bylevel=None,
                                                colsample_bynode=None,
                                                colsample_bytree=None, gamma=None,
                                                gpu_id=None, importance_type='gain',
                                                interaction_constraints=None,
                                                learning_rate=None,
                                                max_delta_step=None, max_depth=None,
                                                min_child_weight=None, missing=nan,
                                                monotone_co...
                                                random_state=42, reg_alpha=None,
                                                reg_lambda=None,
                                                scale_pos_weight=None, subsample=None,
                                                tree_method=None,
                                                validate_parameters=None,
                                                verbosity=None),
                        iid='warn', n_iter=10, n_jobs=None,
```

```
param_distributions={'learning_rate': [0.01, 0.025, 0.05],
                    'n_estimators': [10, 50, 100]},
pre_dispatch='2*n_jobs', random_state=42, refit=True,
return_train_score=False, scoring='neg_mean_absolute_error',
verbose=0)
```

```
[106]: print('The best parameter for', xgb_regressor.best_params_)
print('The best mean absolute error:', (-1) * xgb_regressor.best_score_)
```

```
The best parameter for {'n_estimators': 100, 'learning_rate': 0.05}
The best mean absolute error: 0.11441886818121688
```

Whilst this again reduced the mean absolute error compared to the other tree models, it is still worse than the mean absolute error achieved compared to the ones obtained from the regression models and the tree models have much greater computational cost.

8 Conclusion

In conclusion, the best model to predict SalePrice was the Ridge Regression model which has a mean absolute error of 0.07871660662991475, the lowest of all the models which were trained.

Throughout this notebook, various data science techniques have been used and I have discovered the huge computational cost of using Tree models for very large datasets. The regression models performed well in trying to predict the feature SalePrice.

Now I have completed this project, I am able to evaluate the outcome, more specifically my implementation of the various machine learning techniques used and available to me.

What Went Well:

- Developed many iterative plots and functions which helped speed up exploratory data analysis
- Improved familiarity with Python modules: NumPy, Pandas, SciPy, SkLearn and Seaborn
- Built a machine learning model with an acceptable mean absolute error

Even Better If:

- Used more complex encoding techniques to make better use of categorical features
- used feature engineering to create more useful features to train model on
- Stacked models together and used more hyperparameter tuning techniques

One final thing which would have made my job a lot easier would be to join the training and test files right at the beginning and cleaning both of these datasets together. This would have allowed me to apply the exact same transformations to both datasets. Then I would be able to halve the conjoined dataframe so I could

run separate training and testing sets into my models. But the added benefit would be the processing of both dataframes would be the same.

In further projects I will aim to develop on this improvements I have identified.

I will now store the best model in a submission file. As I log transformed the sale price values, I will reverse this so I retrieve the actual predicted Sale Price values without scaling.

```
[107]: samplesubmission = pd.read_csv(r'sample_submission.csv')

output = pd.DataFrame({'Id': samplesubmission.Id, 'SalePrice': (np.
    ↳exp(rr_pred))})
output.to_csv('submission.csv', index=False)
```